

Final_Project_Product_Search_Relevance_Prediction_

June 25, 2024

#Project Intro & Goals Home and Garden Product Search Relevance Prediction

This project focuses on developing a machine learning model to rate the relevance of search results for home and garden products. Our goal is to create a system that can accurately match search queries with product information, considering factors such as brand, material, and functionality.

To illustrate what we're trying to achieve, let's look at a couple of examples:

1. High Relevance Example: Search term: "angle bracket" Product: Simpson Strong-Tie 12-Gauge Angle Relevance Score: 3.0 (Perfect match)

The product description mentions: "Not only do angles make joints stronger, they also provide more consistent, straight corners. Simpson Strong-Tie offers a wide variety of angles in various sizes and thicknesses..."

This is a perfect match because the search term directly corresponds to the product, and the description confirms its purpose and brand.

2. Lower Relevance Example: Search term: "honda mower" Product: Toro Personal Pace Recycler 22 in. Variable Speed Self-Propelled Gas Lawn Mower with Briggs & Stratton Engine Relevance Score: 2.0 (Partially relevant)

While this product is a lawn mower, it's not a Honda brand. It's partially relevant because it's the right type of product (a mower) but doesn't match the specific brand requested.

Our model aims to predict these relevance scores by analyzing the relationship between search terms, product titles, and detailed product information. We'll be working with a dataset that includes search queries, product titles, descriptions, and attributes. Through data preprocessing, feature engineering, and the application of various machine learning techniques, we aim to build a model that can distinguish between irrelevant, partially relevant, and perfectly matching results.

The ultimate objective is to improve the search experience on e-commerce platforms specializing in home improvement and gardening products by providing more accurate and relevant search results. This notebook will walk through the entire process, from data exploration and cleaning to model development and evaluation, showcasing the steps taken to address this real-world information retrieval challenge.

```
[1]: !pip install Levenshtein
```

Collecting Levenshtein

Downloading

Levenshtein-0.25.1-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl
(177 kB)

177.4/177.4

kB 3.9 MB/s eta 0:00:00

Collecting rapidfuzz<4.0.0,>=3.8.0 (from Levenshtein)

Downloading

rapidfuzz-3.9.3-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (3.4 MB)

3.4/3.4 MB

13.7 MB/s eta 0:00:00

Installing collected packages: rapidfuzz, Levenshtein

Successfully installed Levenshtein-0.25.1 rapidfuzz-3.9.3

```
[2]: import re
import nltk
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from matplotlib_venn import venn2
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from nltk.corpus import stopwords
from nltk.stem.snowball import SnowballStemmer
from nltk.stem import WordNetLemmatizer
import Levenshtein
import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader
from torchvision import transforms
from PIL import Image

# Download necessary NLTK data
nltk.download('wordnet')
nltk.download('stopwords')

# Initialize stemmer, lemmatizer, and stopwords
stemmer = SnowballStemmer('english')
lemmatizer = WordNetLemmatizer()
stop_words = set(stopwords.words('english'))

# Set up matplotlib for inline plotting in notebooks
%matplotlib inline
```

[nltk_data] Downloading package wordnet to /root/nltk_data...

[nltk_data] Downloading package stopwords to /root/nltk_data...

[nltk_data] Unzipping corpora/stopwords.zip.

Here, we load the datasets: train, test, attributes, and product descriptions. This gives us a

comprehensive view of the product information and search queries we'll be working with. The absence of missing values in the train and test sets is a good sign for data quality.

```
[3]: # Load datasets
train_df = pd.read_csv("train.csv", encoding="ISO-8859-1")
test_df = pd.read_csv("test.csv", encoding="ISO-8859-1")
attributes_df = pd.read_csv('attributes.csv', engine='python')
descriptions_df = pd.read_csv('product_descriptions.csv', engine='python')

# Check for missing values
for name, df in [('Train', train_df), ('Test', test_df), ('Attributes', attributes_df), ('Descriptions', descriptions_df)]:
    print(f"\nMissing values in {name} dataset:")
    print(df.isnull().sum())
    print(f"\n{name} Dataset Info:")
    print(df.info())
    print("\nSample data:")
    display(df.head())
    print("\n" + "="*50)
```

Missing values in Train dataset:

```
id          0
product_uid 0
product_title 0
search_term 0
relevance    0
dtype: int64
```

Train Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 74067 entries, 0 to 74066
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id              74067 non-null  int64
1   product_uid     74067 non-null  int64
2   product_title   74067 non-null  object
3   search_term     74067 non-null  object
4   relevance       74067 non-null  float64
dtypes: float64(1), int64(2), object(2)
memory usage: 2.8+ MB
None
```

Sample data:

```
   id  product_uid  product_title \
0   2    100001    Simpson Strong-Tie 12-Gauge Angle
```

1	3	100001	Simpson Strong-Tie 12-Gauge Angle
2	9	100002	BEHR Premium Textured DeckOver 1-gal. #SC-141 ...
3	16	100005	Delta Vero 1-Handle Shower Only Faucet Trim Ki...
4	17	100005	Delta Vero 1-Handle Shower Only Faucet Trim Ki...

	search_term	relevance
0	angle bracket	3.00
1	l bracket	2.50
2	deck over	3.00
3	rain shower head	2.33
4	shower only faucet	2.67

=====

Missing values in Test dataset:

```
id          0
product_uid 0
product_title 0
search_term 0
dtype: int64
```

Test Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 166693 entries, 0 to 166692

Data columns (total 4 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	id	166693 non-null	int64
1	product_uid	166693 non-null	int64
2	product_title	166693 non-null	object
3	search_term	166693 non-null	object

dtypes: int64(2), object(2)

memory usage: 5.1+ MB

None

Sample data:

	id	product_uid	product_title \
0	1	100001	Simpson Strong-Tie 12-Gauge Angle
1	4	100001	Simpson Strong-Tie 12-Gauge Angle
2	5	100001	Simpson Strong-Tie 12-Gauge Angle
3	6	100001	Simpson Strong-Tie 12-Gauge Angle
4	7	100001	Simpson Strong-Tie 12-Gauge Angle

	search_term
0	90 degree bracket
1	metal l brackets
2	simpson sku able

```

3      simpson strong  ties
4  simpson strong tie hcc668

```

=====

Missing values in Attributes dataset:

```

product_uid      155
name             155
value           6422
dtype: int64

```

Attributes Dataset Info:

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2044803 entries, 0 to 2044802
Data columns (total 3 columns):
#   Column      Dtype
---  -
0   product_uid  float64
1   name         object
2   value        object
dtypes: float64(1), object(2)
memory usage: 46.8+ MB
None

```

Sample data:

	product_uid	name	value
0	100001.0	Bullet01	Versatile connector for various 90° connection...
1	100001.0	Bullet02	Stronger than angled nailing or screw fastenin...
2	100001.0	Bullet03	Help ensure joints are consistently straight a...
3	100001.0	Bullet04	Dimensions: 3 in. x 3 in. x 1-1/2 in.
4	100001.0	Bullet05	Made from 12-Gauge steel

=====

Missing values in Descriptions dataset:

```

product_uid      0
product_description  0
dtype: int64

```

Descriptions Dataset Info:

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 124428 entries, 0 to 124427
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   product_uid  124428 non-null  int64

```

```
1    product_description    124428 non-null    object
dtypes: int64(1), object(1)
memory usage: 1.9+ MB
None
```

Sample data:

	product_uid	product_description
0	100001	Not only do angles make joints stronger, they ...
1	100002	BEHR Premium Textured DECKOVER is an innovativ...
2	100003	Classic architecture meets contemporary design...
3	100004	The Grape Solar 265-Watt Polycrystalline PV So...
4	100005	Update your bathroom with the Delta Vero Singl...

=====

Dataset sizes:

Train dataset: 74,067 rows

Test dataset: 166,693 rows

Attributes dataset: 2,044,803 rows

Descriptions dataset: 124,428 rows

Original features:

Train dataset: 5 features (id, product_uid, product_title, search_term, relevance)

Test dataset: 4 features (id, product_uid, product_title, search_term)

#Visualizing the data:

```
[4]: # Analyze relevance distribution in train dataset
plt.figure(figsize=(10, 6))
sns.histplot(data=train_df, x='relevance', bins=13, kde=True)
plt.title('Distribution of Relevance Scores', fontsize=16)
plt.xlabel('Relevance Score', fontsize=14)
plt.ylabel('Frequency', fontsize=14)
plt.grid(True, linestyle='--', alpha=0.7)
plt.show()


# Compare search terms between train and test sets
train_search_terms = set(train_df['search_term'])
test_search_terms = set(test_df['search_term'])
venn2([train_search_terms, test_search_terms], ('Train search terms', 'Test_
↪search terms'))
plt.title('Overlap of Search Terms in Train and Test Sets')
plt.show()

# Compare product_uid between train and test sets
```


```

train_product_uids = set(train_df['product_uid'])
test_product_uids = set(test_df['product_uid'])
venn2([train_product_uids, test_product_uids], ('Train product_uid', 'Test_
↪product_uid'))
plt.title('Overlap of Product UIDs in Train and Test Sets')
plt.show()

```



Overlap of Search Terms in Train and Test Sets



Overlap of Product UIDs in Train and Test Sets



This histogram shows the distribution of relevance scores in the training data. We can see that the scores are not evenly distributed:

There's a peak at 3.0, indicating many perfect matches. Another peak around 2.3-2.7, suggesting many partially relevant results. Fewer scores below 2.0, implying that completely irrelevant results are less common. This distribution suggests that the search algorithm is generally performing well, but there's room for improvement in distinguishing between partial and perfect matches.

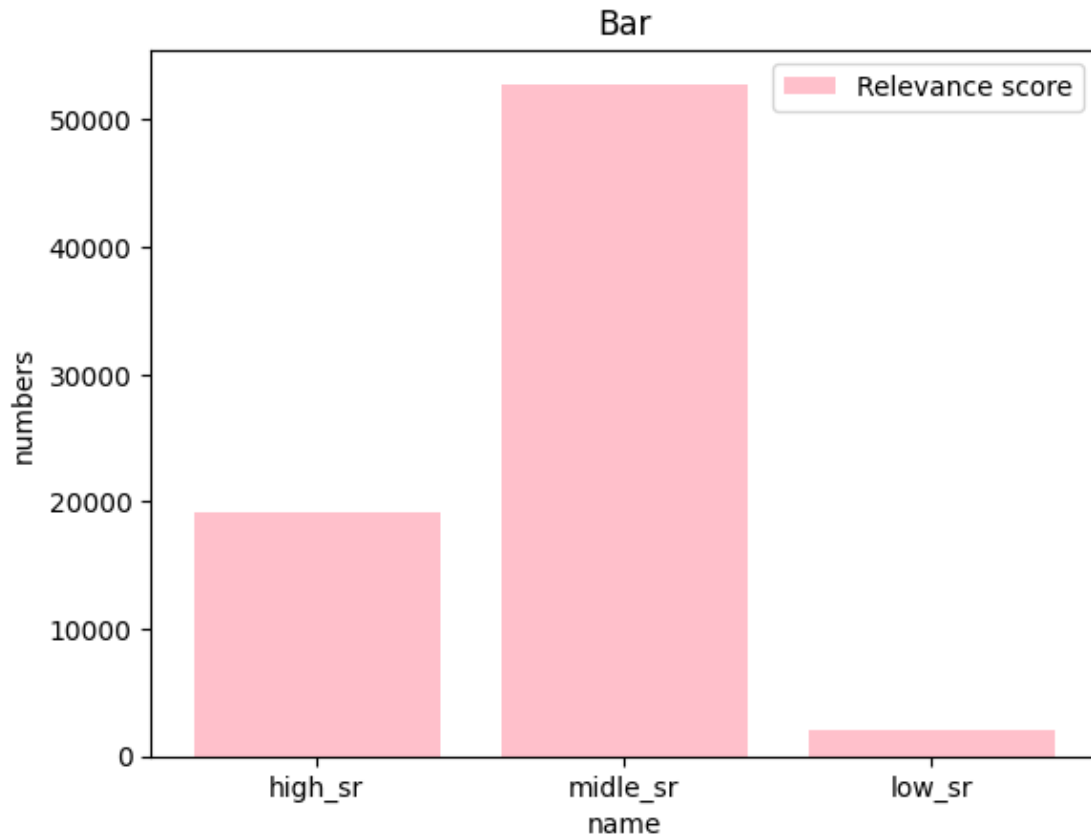
These Venn diagrams show the overlap between train and test sets:

1. Search terms: There's significant overlap, but also many unique terms in each set. This suggests that our model needs to generalize well to new search queries.

2. Product UUIDs: Almost all products in the train set are also in the test set, with many additional products in the test set. This indicates that our model will be evaluated on both familiar and new products.

```
[38]: def count_SR():  
    high_SR = [relevance for relevance in train_df['relevance'] if relevance ==  
    ↪ 3.00]  
    middle_SR = [relevance for relevance in train_df['relevance'] if 1.00 <  
    ↪ relevance < 3.00]  
    low_SR = [relevance for relevance in train_df['relevance'] if relevance ==  
    ↪ 1.00]  
    return [len(high_SR), len(middle_SR), len(low_SR)]
```

```
[6]: plt.bar(['high_sr', 'midle_sr', 'low_sr'], count_SR(), label='Relevance score',  
    ↪ color='pink')  
plt.xlabel('name')  
plt.ylabel('numbers')  
plt.title('Bar')  
plt.legend()  
plt.show()
```



These cells provide more detailed breakdowns of the relevance scores:

The bar chart shows the count of high, middle, and low relevance scores. The count plot confirms the distribution we saw in the histogram.


These visualizations reinforce that partially relevant results are the most common, followed by perfect matches, with fewer irrelevant results.

```
[7]: sns.countplot(x="relevance", data=train_df, palette="Reds")  
plt.show()
```

<ipython-input-7-8f793203aa76>:1: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x="relevance", data=train_df, palette="Reds")
```



#Cleaning the data:

This section is crucial for preparing our data for machine learning:

We clean the text data by removing special characters, standardizing abbreviations, and applying stemming and lemmatization.

We merge product descriptions with attributes to create a comprehensive product representation.

We handle missing values to ensure data completeness. These steps are essential for creating meaningful features that capture the relationship between search queries and product information.

```
[8]: def remove_duplicates(string):
    unique_tokens = []
    [unique_tokens.append(str(token)) for token in string.split() if token not
    ↪ in unique_tokens]
    return ' '.join(unique_tokens)

def clean_text(text):
    text = str(text).lower()

    # Replace common abbreviations and remove specific strings
    replacements = {
```

```

        'in.': 'inch', 'ft.': 'foot', '-oz.': 'ounce', 'oz.': 'ounce',
        'sq.': 'square', 'gal.': 'gallon', 'lb.': 'pound', 'cu.': 'cubic',
        'o.d.': 'outer diameter', 'dia.': 'diameter', '-': '', 'r ': 'r'
    }
    for old, new in replacements.items():
        text = text.replace(old, new)

    # Remove URLs and specific phrases
    text = re.sub(r'http\S+', '', text)
    text = re.sub(r'click here.*', '', text)
    text = re.sub(r'please visit.*', '', text)

    # Remove special characters and extra spaces
    text = re.sub(r'[^a-zA-Z0-9]+', ' ', text)
    text = re.sub(r'\s+', ' ', text).strip()

    return text

def preprocess_text(text):
    # Clean the text
    text = clean_text(text)

    # Remove stopwords
    words = [word for word in text.split() if word not in stop_words]

    # Lemmatize and stem
    words = [lemmatizer.lemmatize(stemmer.stem(word)) for word in words]

    return ' '.join(words)

# Define strings to remove
strings_to_remove = [
    'br', 'src', 'href', 'alt', 'please visit',
    'Click here to review our return policy for additional information_
    ↪regarding returns',
    'Click here to see Home Depot',
    'Click here for our Project Guide',
    'Click here for our Buying Guide',
    'Click on the More Info tab to download',
    'CLICK HERE to create your own collection',
    'Click Here for details on the services',
    'Click Here for Ideas and Designs',
    'Click Here for a Demo of the Design',
    'Click Here to learn more about',
    'CLICK HERE to view our',
    'Click below to visit our',
    'Click here to purchase a sample of this',

```

```

        'click on the link to get started',
        'Click image to enlarge',
        'https://www.ryobitools.com/nation',
        'http://www.homedepot.com/ApplianceDeliveryandInstallation',
        'http://itemvideo-dev.microsite.homedepot.com/111414/26P/
↳online_BB_banner_111114.jpg',
        'http://www.homedepot.com/p/Rev-A-Shelf-Door-Mounting-Kit-5WB-DMKIT/
↳202855698'
    ]

```

Handle missing values and process attributes and descriptions:

```

[9]: # Handle missing values in attributes
attributes_df['value'].fillna('', inplace=True)

# Process attributes
attributes_df['name'] = attributes_df['name'].apply(lambda x: x[:6] if 'bullet' in
↳str(x).lower() else x)
attributes_df['product_attributes'] = attributes_df['name'] + ' ' +
↳attributes_df['value']
attributes_df = attributes_df.groupby('product_uid')['product_attributes'].
↳apply(' '.join).reset_index()
attributes_df['clean_attributes'] = attributes_df['product_attributes'].
↳apply(preprocess_text)
attributes_df['clean_attributes'] = attributes_df['clean_attributes'].
↳apply(remove_duplicates)

# Process descriptions
for string in strings_to_remove:
    descriptions_df['product_description'] =
↳descriptions_df['product_description'].apply(lambda x: x.lower().
↳replace(string.lower(), ''))

descriptions_df['clean_description'] = descriptions_df['product_description'].
↳apply(preprocess_text)

# Merge descriptions and attributes
df_des_attr = pd.merge(descriptions_df, attributes_df, on='product_uid',
↳how='left')
df_des_attr['product_description_attributes'] =
↳df_des_attr['clean_description'] + ' ' + df_des_attr['clean_attributes']

# Handle missing values in merged dataframe
df_des_attr['product_description_attributes'].fillna('', inplace=True)

print("Missing values after processing:")

```

```
print(df_des_attr.isnull().sum())
```

Missing values after processing:

```
product_uid          0
product_description   0
clean_description     0
product_attributes    38165
clean_attributes      38165
product_description_attributes  0
dtype: int64
```

```
[10]: df_des_attr.head(5)
```

```
[10]:  product_uid          product_description \
0      100001  not only do angles make joints stronger, they ...
1      100002  behr premium textured deckover is an innovativ...
2      100003  classic architecture meets contemporary design...
3      100004  the grape solar 265-watt polycrystalline pv so...
4      100005  update your bathroom with the delta vero singl...

          clean_description \
0  angl make joint stronger also provid consist s...
1  behrpremium textur deckoveri innov solid color...
2  classic architectur meet contemporari design e...
3  grape solar265watt polycrystallin pv solarpane...
4  updat yourbathroom delta vero singlehandl show...

          product_attributes \
0  Bullet Versatile connector for various 90° con...
1  Application Method Brush,Roller,Spray Assemble...
2  Built-in flange Yes Bullet Slightly narrower f...
3  Amperage (amps) 8.56 Bullet Positive power tol...
4  Bath Faucet Type Combo Tub and Shower Built-in...

          clean_attributes \
0  bullet versatil connectorforvari 90 connect ho...
1  applic method brush roller spray assembl depth...
2  builtin flang yes bullet slight narrowerfortig...
3  amperag amp 8 56 bullet posit powertoler 0 5wa...
4  bath faucet type combo tub showerbuiltin water...

          product_description_attributes
0  angl make joint stronger also provid consist s...
1  behrpremium textur deckoveri innov solid color...
2  classic architectur meet contemporari design e...
3  grape solar265watt polycrystallin pv solarpane...
4  updat yourbathroom delta vero singlehandl show...
```

Process train and test datasets:

```
[11]: # Function to process a single dataframe
def process_dataframe(df, df_des_attr):
    df['clean_title'] = df['product_title'].apply(preprocess_text)
    df['clean_search_term'] = df['search_term'].apply(preprocess_text)

    # Remove any existing 'product_description_attributes' column
    if 'product_description_attributes' in df.columns:
        df = df.drop('product_description_attributes', axis=1)

    # Merge with df_des_attr
    merged_df = pd.merge(df, df_des_attr[['product_uid',
    ↪ 'product_description_attributes']], on='product_uid', how='left')

    # Handle any missing values after merging
    merged_df['product_description_attributes'].fillna('', inplace=True)

    return merged_df

# Process train and test datasets
train_df = process_dataframe(train_df, df_des_attr)
test_df = process_dataframe(test_df, df_des_attr)

# Remove duplicate columns if they exist
for df in [train_df, test_df]:
    columns_to_drop = [col for col in df.columns if col.
    ↪ startswith('product_description_attributes_')]
    df.drop(columns=columns_to_drop, inplace=True)

# Print information about the processed datasets
for name, df in [('Train', train_df), ('Test', test_df)]:
    print(f"\nProcessing of {name} dataset completed.")
    print(f"Columns in {name} dataset:")
    print(df.columns)
    print(f"\nMissing values in processed {name} dataset:")
    print(df.isnull().sum())
    print(f"\nSample of processed {name} data:")
    display(df[['product_uid', 'clean_title', 'clean_search_term',
    ↪ 'product_description_attributes']].head())

print("\nData preprocessing completed.")

# Final verification
print("\nFinal Train dataset columns:")
print(train_df.columns)
print("\nFinal Test dataset columns:")
```

```

print(test_df.columns)

# Check if the column was added successfully
if 'product_description_attributes' in train_df.columns and
    'product_description_attributes' in test_df.columns:
    print("\nThe 'product_description_attributes' column is present in both
    datasets.")
else:
    print("\nWarning: The 'product_description_attributes' column is missing
    from one or both datasets.")
    print("Train columns:", 'product_description_attributes' in train_df.
    columns)
    print("Test columns:", 'product_description_attributes' in test_df.columns)

# Additional check: print the number of rows in each dataset before and after
merging
print("\nNumber of rows:")
print(f"Original train_df: {len(train_df)}")
print(f"Original test_df: {len(test_df)}")
print(f"df_des_attr: {len(df_des_attr)}")

```

Processing of Train dataset completed.

Columns in Train dataset:

```

Index(['id', 'product_uid', 'product_title', 'search_term', 'relevance',
       'clean_title', 'clean_search_term', 'product_description_attributes'],
      dtype='object')

```

Missing values in processed Train dataset:

```

id                0
product_uid       0
product_title     0
search_term       0
relevance         0
clean_title       0
clean_search_term 0
product_description_attributes 0
dtype: int64

```

Sample of processed Train data:

```

      product_uid      clean_title \
0      100001      simpson strongti 12gaug angl
1      100001      simpson strongti 12gaug angl
2      100002  behrpremium textur deckover1gallon sc141 tugbo...
3      100005  delta vero 1handl showeron faucet trim kit chr...
4      100005  delta vero 1handl showeron faucet trim kit chr...

```


	clean_search_term	product_description_attributes
0	angl bracket	angl make joint stronger also provid consist s...
1	l bracket	angl make joint stronger also provid consist s...
2	deck	behrpremium textur deckoveri innov solid color...
3	rain showerhead	updat yourbathroom delta vero singlehandl show...
4	showeron faucet	updat yourbathroom delta vero singlehandl show...

Processing of Test dataset completed.

Columns in Test dataset:

```
Index(['id', 'product_uid', 'product_title', 'search_term', 'clean_title',
      'clean_search_term', 'product_description_attributes'],
      dtype='object')
```

Missing values in processed Test dataset:

```
id                0
product_uid       0
product_title     0
search_term       0
clean_title       0
clean_search_term 0
product_description_attributes 0
dtype: int64
```

Sample of processed Test data:

	product_uid	clean_title	clean_search_term \
0	100001	simpson strongti 12gaug angl	90 degre bracket
1	100001	simpson strongti 12gaug angl	metal l bracket
2	100001	simpson strongti 12gaug angl	simpson sku abl
3	100001	simpson strongti 12gaug angl	simpson strong tie
4	100001	simpson strongti 12gaug angl	simpson strong tie hcc668

	product_description_attributes
0	angl make joint stronger also provid consist s...
1	angl make joint stronger also provid consist s...
2	angl make joint stronger also provid consist s...
3	angl make joint stronger also provid consist s...
4	angl make joint stronger also provid consist s...

Data preprocessing completed.

Final Train dataset columns:

```
Index(['id', 'product_uid', 'product_title', 'search_term', 'relevance',
      'clean_title', 'clean_search_term', 'product_description_attributes'],
      dtype='object')
```

Final Test dataset columns:

```
Index(['id', 'product_uid', 'product_title', 'search_term', 'clean_title',  
      'clean_search_term', 'product_description_attributes'],  
      dtype='object')
```

The 'product_description_attributes' column is present in both datasets.

Number of rows:

Original train_df: 74067

Original test_df: 166693

df_des_attr: 124428

After initial preprocessing:

Train dataset: 8 features

Test dataset: 7 features

Feature engineering

Here we create several important features:

1. Levenshtein ratio: Measures string similarity between search term and product title.
2. Title length, search length, and description length: Capture the amount of information available.
3. Various 'shared words' features: Measure overlap between search terms and product information.

These features aim to quantify different aspects of the match between search queries and products, which will be crucial for predicting relevance.

```
[12]: import numpy as np  
  
# Levenshtein distance between search term and product title  
def levenshtein_ratio(row):  
    try:  
        return Levenshtein.ratio(str(row['clean_search_term']),  
                                   str(row['clean_title']))  
    except:  
        return np.nan  
  
for name, df in [('Train', train_df), ('Test', test_df)]:  
    print(f"\nProcessing {name} dataset:")  
  
    df['levenshtein_ratio'] = df.apply(levenshtein_ratio, axis=1)  
    df['title_len'] = df['clean_title'].apply(len)  
    df['search_len'] = df['clean_search_term'].apply(len)  
    df['description_len'] = df['product_description_attributes'].apply(len)  
  
    print(f"Feature engineering completed for {name} dataset.")  
    print(f"Sample of {name} data with engineered features:")
```

```

display(df[['clean_search_term', 'clean_title', 'levenshtein_ratio',
↪ 'title_len', 'search_len', 'description_len']].head())

print(f"\n{name} dataset structure:")
print(df.info())

print(f"\nMissing values in {name} dataset:")
print(df.isnull().sum())

print(f"\nDescriptive statistics for {name} dataset:")
print(df[['levenshtein_ratio', 'title_len', 'search_len',
↪ 'description_len']].describe())

print("\nFeature engineering completed for both datasets.")

# Correlation matrix for numerical features in train dataset
correlation_matrix = train_df[['levenshtein_ratio', 'title_len', 'search_len',
↪ 'description_len', 'relevance']].corr()
print("\nCorrelation matrix for numerical features in train dataset:")
print(correlation_matrix)

# Visualize correlation matrix
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', vmin=-1, vmax=1,
↪ center=0)
plt.title('Correlation Heatmap of Numerical Features')
plt.show()

```

Processing Train dataset:

Feature engineering completed for Train dataset.

Sample of Train data with engineered features:

	clean_search_term	clean_title \
0	angl bracket	simpson strongti 12gaug angl
1	l bracket	simpson strongti 12gaug angl
2	deck behrpremium textur deckover1gallon sc141 tugbo...	
3	rain showerhead delta vero 1handl showeron faucet trim kit chr...	
4	showeron faucet delta vero 1handl showeron faucet trim kit chr...	

	levenshtein_ratio	title_len	search_len	description_len
0	0.200000	28	12	970
1	0.162162	28	9	970
2	0.114286	66	4	1634
3	0.368421	61	15	1073

4	0.394737	61	15	1073
---	----------	----	----	------

Train dataset structure:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 74067 entries, 0 to 74066

Data columns (total 12 columns):

#	Column	Non-Null Count	Dtype
0	id	74067 non-null	int64
1	product_uid	74067 non-null	int64
2	product_title	74067 non-null	object
3	search_term	74067 non-null	object
4	relevance	74067 non-null	float64
5	clean_title	74067 non-null	object
6	clean_search_term	74067 non-null	object
7	product_description_attributes	74067 non-null	object
8	levenshtein_ratio	74067 non-null	float64
9	title_len	74067 non-null	int64
10	search_len	74067 non-null	int64
11	description_len	74067 non-null	int64

dtypes: float64(2), int64(5), object(5)

memory usage: 6.8+ MB

None

Missing values in Train dataset:

id	0
product_uid	0
product_title	0
search_term	0
relevance	0
clean_title	0
clean_search_term	0
product_description_attributes	0
levenshtein_ratio	0
title_len	0
search_len	0
description_len	0

dtype: int64

Descriptive statistics for Train dataset:

	levenshtein_ratio	title_len	search_len	description_len
count	74067.000000	74067.000000	74067.000000	74067.000000
mean	0.340163	62.315971	16.906517	913.496145
std	0.120614	20.716386	6.872335	738.723762
min	0.000000	8.000000	0.000000	0.000000
25%	0.256410	47.000000	12.000000	474.000000
50%	0.327869	60.000000	16.000000	869.000000

75%	0.412371	76.000000	21.000000	1293.000000
max	1.000000	144.000000	61.000000	6664.000000

Processing Test dataset:

Feature engineering completed for Test dataset.

Sample of Test data with engineered features:

	clean_search_term		clean_title	levenshtein_ratio	\
0	90 degree bracket	simpson strongti	12gaug angl	0.181818	
1	metal l bracket	simpson strongti	12gaug angl	0.232558	
2	simpson sku abl	simpson strongti	12gaug angl	0.604651	
3	simpson strong tie	simpson strongti	12gaug angl	0.695652	
4	simpson strong tie hcc668	simpson strongti	12gaug angl	0.641509	

	title_len	search_len	description_len
0	28	16	970
1	28	15	970
2	28	15	970
3	28	18	970
4	28	25	970

Test dataset structure:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 166693 entries, 0 to 166692

Data columns (total 11 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	id	166693 non-null	int64
1	product_uid	166693 non-null	int64
2	product_title	166693 non-null	object
3	search_term	166693 non-null	object
4	clean_title	166693 non-null	object
5	clean_search_term	166693 non-null	object
6	product_description_attributes	166693 non-null	object
7	levenshtein_ratio	166693 non-null	float64
8	title_len	166693 non-null	int64
9	search_len	166693 non-null	int64
10	description_len	166693 non-null	int64

dtypes: float64(1), int64(5), object(5)

memory usage: 14.0+ MB

None

Missing values in Test dataset:

id	0
product_uid	0
product_title	0
search_term	0
clean_title	0

```

clean_search_term          0
product_description_attributes  0
levenshtein_ratio          0
title_len                  0
search_len                 0
description_len             0
dtype: int64

```

Descriptive statistics for Test dataset:

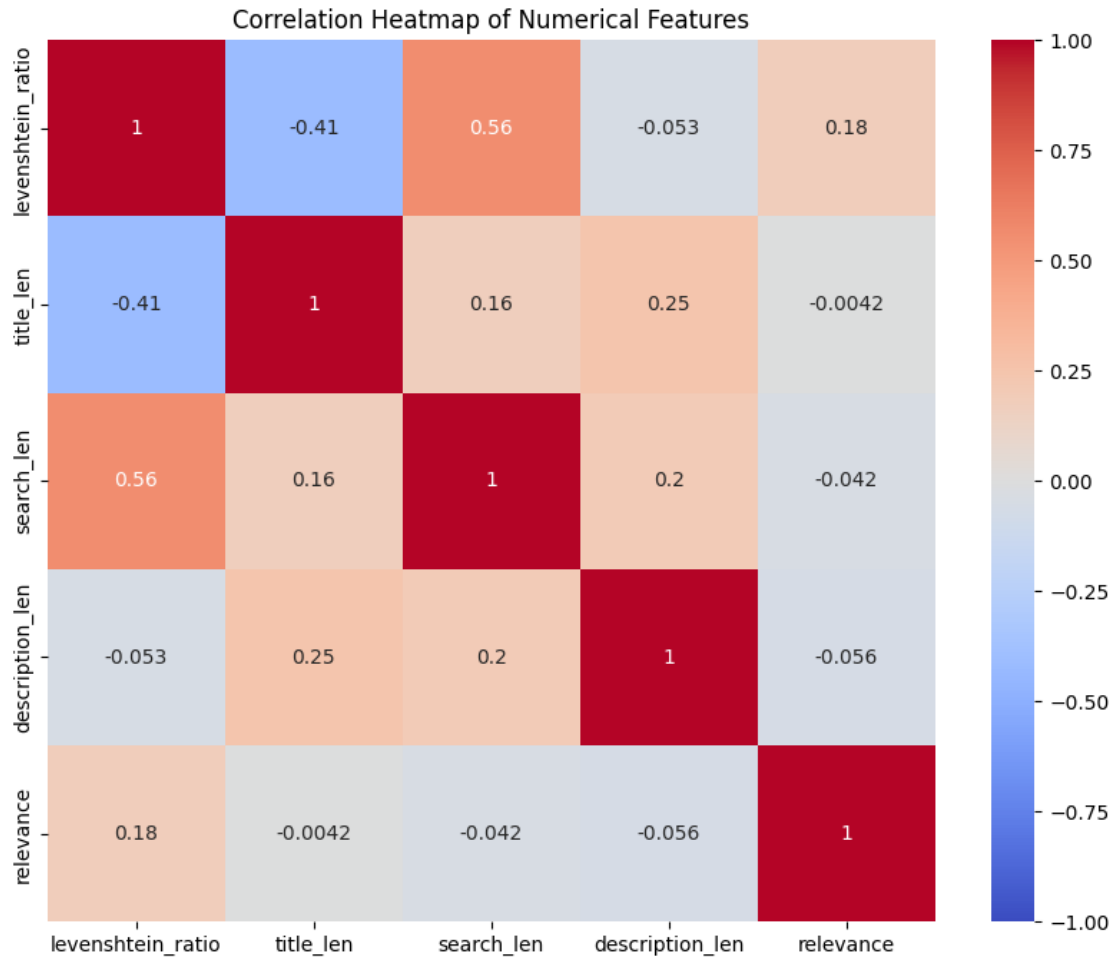
	levenshtein_ratio	title_len	search_len	description_len
count	166693.000000	166693.000000	166693.000000	166693.000000
mean	0.332418	61.912798	15.886432	988.582838
std	0.121593	20.849197	6.765947	710.646165
min	0.000000	7.000000	0.000000	0.000000
25%	0.246575	46.000000	11.000000	605.000000
50%	0.320000	60.000000	15.000000	918.000000
75%	0.405405	75.000000	20.000000	1340.000000
max	1.000000	144.000000	61.000000	6761.000000

Feature engineering completed for both datasets.

Correlation matrix for numerical features in train dataset:

	levenshtein_ratio	title_len	search_len	description_len	\
levenshtein_ratio	1.000000	-0.414973	0.556353	-0.052651	
title_len	-0.414973	1.000000	0.164286	0.248520	
search_len	0.556353	0.164286	1.000000	0.196947	
description_len	-0.052651	0.248520	0.196947	1.000000	
relevance	0.179064	-0.004172	-0.041905	-0.055704	

	relevance
levenshtein_ratio	0.179064
title_len	-0.004172
search_len	-0.041905
description_len	-0.055704
relevance	1.000000



##After feature engineering:

Train dataset: 12 features

Test dataset: 11 features

New features added: - levenshtein_ratio - title_len - search_len - description_len

#Correlation Analysis and Final Data Preparation

[13]: *# Cell 9: Handling Outliers and Final Data Preparation*

```
import numpy as np
import pandas as pd
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import RobustScaler

# Function to apply winsorization
def winsorize(data, limits=(0.05, 0.05)):
```

```

    lower = np.percentile(data, limits[0] * 100)
    upper = np.percentile(data, (1 - limits[1]) * 100)
    return np.clip(data, lower, upper)

# Function to check for outliers
def check_outliers(data, name):
    Q1 = data.quantile(0.25)
    Q3 = data.quantile(0.75)
    IQR = Q3 - Q1
    outliers = ((data < (Q1 - 1.5 * IQR)) | (data > (Q3 + 1.5 * IQR))).sum()
    print(f"\nOutliers in {name} data:")
    print(outliers)
    return outliers.sum()

# Prepare features for modeling
features = ['levenshtein_ratio', 'title_len', 'search_len', 'description_len']
X_train = train_df[features]
y_train = train_df['relevance']
X_test = test_df[features]

# Original data statistics and outlier check
print("Original data statistics:")
print(X_train.describe())
original_outliers = check_outliers(X_train, "original")

# 1. Winsorization
X_train_winsorized = X_train.copy()
X_test_winsorized = X_test.copy()
for column in features:
    X_train_winsorized[column] = winsorize(X_train[column])
    X_test_winsorized[column] = winsorize(X_test[column])

print("\nWinsorized data statistics:")
print(X_train_winsorized.describe())
winsorized_outliers = check_outliers(X_train_winsorized, "winsorized")

# 2. Log transformation
X_train_log = X_train.copy()
X_test_log = X_test.copy()
for column in features:
    X_train_log[column] = np.log1p(X_train[column])
    X_test_log[column] = np.log1p(X_test[column])

print("\nLog-transformed data statistics:")
print(X_train_log.describe())
log_outliers = check_outliers(X_train_log, "log-transformed")

```



```

# 3. Robust scaling
scaler = RobustScaler()
X_train_robust = pd.DataFrame(scaler.fit_transform(X_train), columns=features)
X_test_robust = pd.DataFrame(scaler.transform(X_test), columns=features)

print("\nRobust scaled data statistics:")
print(X_train_robust.describe())
robust_outliers = check_outliers(X_train_robust, "robust scaled")

# Check for any remaining missing values
for name, data in [("Original", X_train), ("Winsorized", X_train_winsorized),
                  ("Log-transformed", X_train_log), ("Robust scaled", X_train_robust)]:
    print(f"\nMissing values in {name} X_train:")
    print(data.isnull().sum())

print("\nTotal number of outliers:")
print(f"Original: {original_outliers}")
print(f"Winsorized: {winsorized_outliers}")
print(f"Log-transformed: {log_outliers}")
print(f"Robust scaled: {robust_outliers}")

print("\nData is now ready for modeling with different outlier handling methods.")

# You can choose which version of the data to use for your modeling
# For example:
# X_train_final = X_train_winsorized
# X_test_final = X_test_winsorized

```

Original data statistics:

	levenshtein_ratio	title_len	search_len	description_len
count	74067.000000	74067.000000	74067.000000	74067.000000
mean	0.340163	62.315971	16.906517	913.496145
std	0.120614	20.716386	6.872335	738.723762
min	0.000000	8.000000	0.000000	0.000000
25%	0.256410	47.000000	12.000000	474.000000
50%	0.327869	60.000000	16.000000	869.000000
75%	0.412371	76.000000	21.000000	1293.000000
max	1.000000	144.000000	61.000000	6664.000000

Outliers in original data:

levenshtein_ratio	1166
title_len	503
search_len	1194
description_len	2051
dtype:	int64

Winsorized data statistics:

	levenshtein_ratio	title_len	search_len	description_len
count	74067.000000	74067.000000	74067.000000	74067.000000
mean	0.338331	62.065859	16.761135	883.141264
std	0.108102	19.115123	6.100321	652.159238
min	0.163636	32.000000	7.000000	0.000000
25%	0.256410	47.000000	12.000000	474.000000
50%	0.327869	60.000000	16.000000	869.000000
75%	0.412371	76.000000	21.000000	1293.000000
max	0.556962	100.000000	29.000000	2214.000000

Outliers in winsorized data:

```
levenshtein_ratio    0
title_len            0
search_len           0
description_len       0
dtype: int64
```

Log-transformed data statistics:

	levenshtein_ratio	title_len	search_len	description_len
count	74067.000000	74067.000000	74067.000000	74067.000000
mean	0.288832	4.092331	2.806722	5.329620
std	0.088575	0.341882	0.412514	2.992691
min	0.000000	2.197225	0.000000	0.000000
25%	0.228259	3.871201	2.564949	6.163315
50%	0.283575	4.110874	2.833213	6.768493
75%	0.345270	4.343805	3.091042	7.165493
max	0.693147	4.976734	4.127134	8.804625

Outliers in log-transformed data:

```
levenshtein_ratio    781
title_len            498
search_len           1233
description_len       17544
dtype: int64
```

Robust scaled data statistics:

	levenshtein_ratio	title_len	search_len	description_len
count	74067.000000	74067.000000	74067.000000	74067.000000
mean	0.078829	0.079861	0.100724	0.054330
std	0.773363	0.714358	0.763593	0.901983
min	-2.102251	-1.793103	-1.777778	-1.061050
25%	-0.458183	-0.448276	-0.444444	-0.482295
50%	0.000000	0.000000	0.000000	0.000000
75%	0.541817	0.551724	0.555556	0.517705
max	4.309614	2.896552	5.000000	7.075702

Outliers in robust scaled data:

```
levenshtein_ratio    1166
title_len            503
search_len           1194
description_len       2051
dtype: int64
```

Missing values in Original X_train:

```
levenshtein_ratio    0
title_len            0
search_len           0
description_len       0
dtype: int64
```

Missing values in Winsorized X_train:

```
levenshtein_ratio    0
title_len            0
search_len           0
description_len       0
dtype: int64
```

Missing values in Log-transformed X_train:

```
levenshtein_ratio    0
title_len            0
search_len           0
description_len       0
dtype: int64
```

Missing values in Robust scaled X_train:

```
levenshtein_ratio    0
title_len            0
search_len           0
description_len       0
dtype: int64
```

Total number of outliers:

```
Original: 4914
Winsorized: 0
Log-transformed: 20056
Robust scaled: 4914
```

Data is now ready for modeling with different outlier handling methods.

#Feature Analysis

Import Libraries and Setup:

```
[14]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_selection import mutual_info_regression
from sklearn.ensemble import RandomForestRegressor
from scipy.stats import spearmanr
from mpl_toolkits.mplot3d import Axes3D

plt.style.use('seaborn')
sns.set_palette("deep")
```

<ipython-input-14-320d71af9d2f>:10: MatplotlibDeprecationWarning: The seaborn styles shipped by Matplotlib are deprecated since 3.6, as they no longer correspond to the styles shipped by seaborn. However, they will remain available as 'seaborn-v0_8-*<style>*'. Alternatively, directly use the seaborn API instead.

```
plt.style.use('seaborn')
```

Helper Functions

```
[15]: def plot_feature_distribution(df, feature, title):
    plt.figure(figsize=(10, 6))
    sns.histplot(df[feature], kde=True)
    plt.title(f'Distribution of {title}')
    plt.xlabel(feature)
    plt.ylabel('Count')
    plt.show()

def plot_feature_vs_relevance(df, feature, title):
    plt.figure(figsize=(10, 6))
    sns.scatterplot(x=feature, y='relevance', data=df, alpha=0.5)
    plt.title(f'{title} vs Relevance')
    plt.xlabel(feature)
    plt.ylabel('Relevance')


    # Add regression line
    x = df[feature].values.reshape(-1, 1)
    y = df['relevance'].values
    plt.plot(np.unique(x), np.poly1d(np.polyfit(x.ravel(), y, 1))(np.
↪unique(x)), color='red')

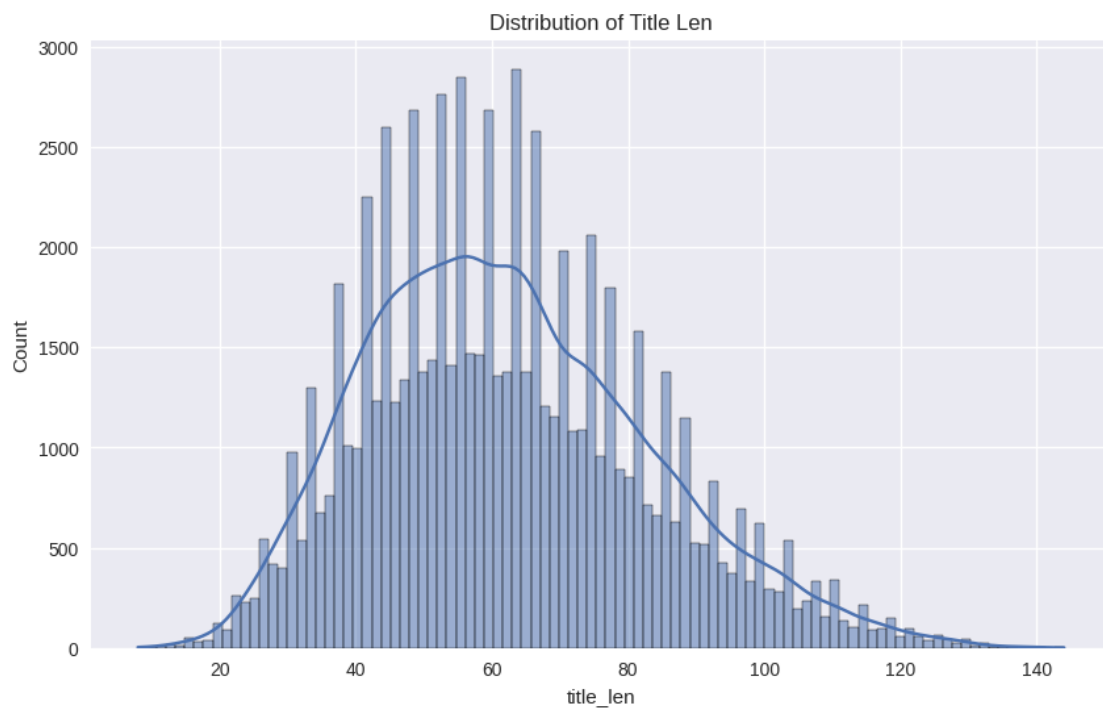
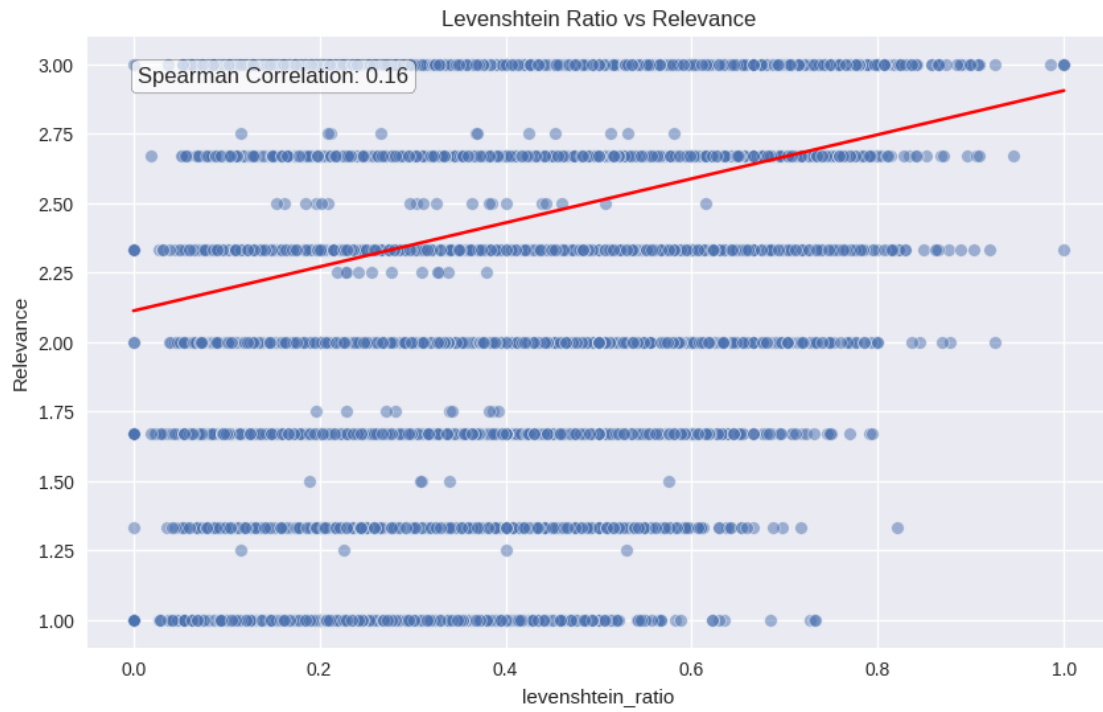
    # Add Spearman correlation
    corr, _ = spearmanr(df[feature], df['relevance'])
    plt.text(0.05, 0.95, f'Spearman Correlation: {corr:.2f}', transform=plt.
↪gca().transAxes,
            verticalalignment='top', fontsize=12, bbox=dict(boxstyle='round',
↪facecolor='white', alpha=0.7))
```

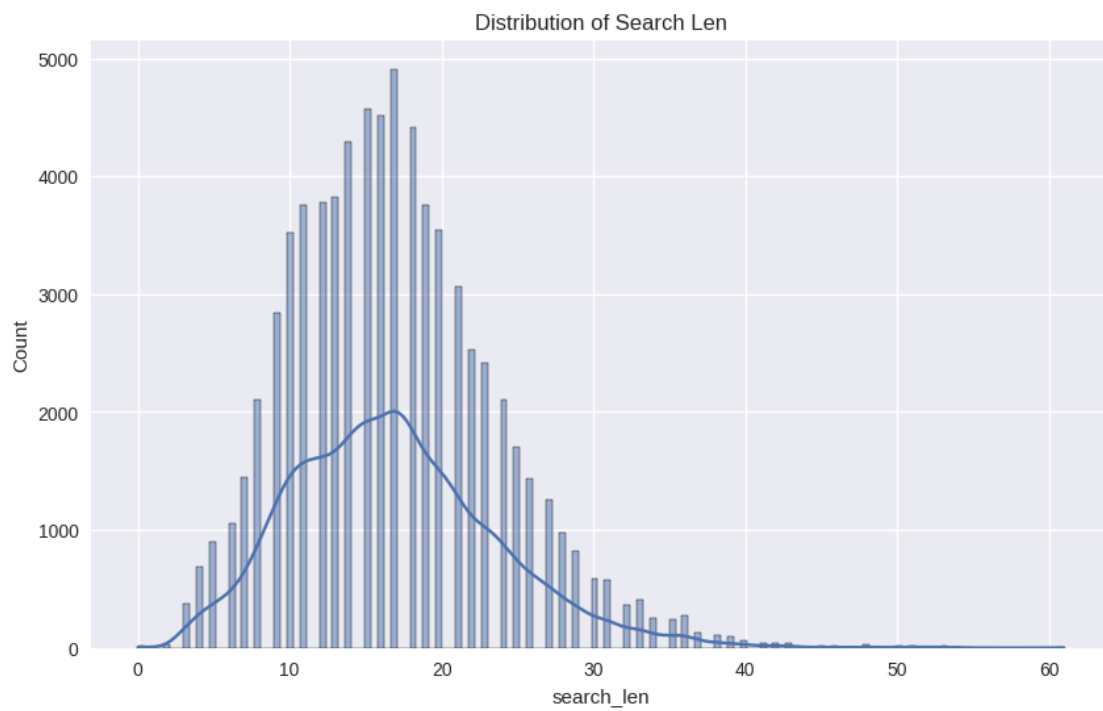
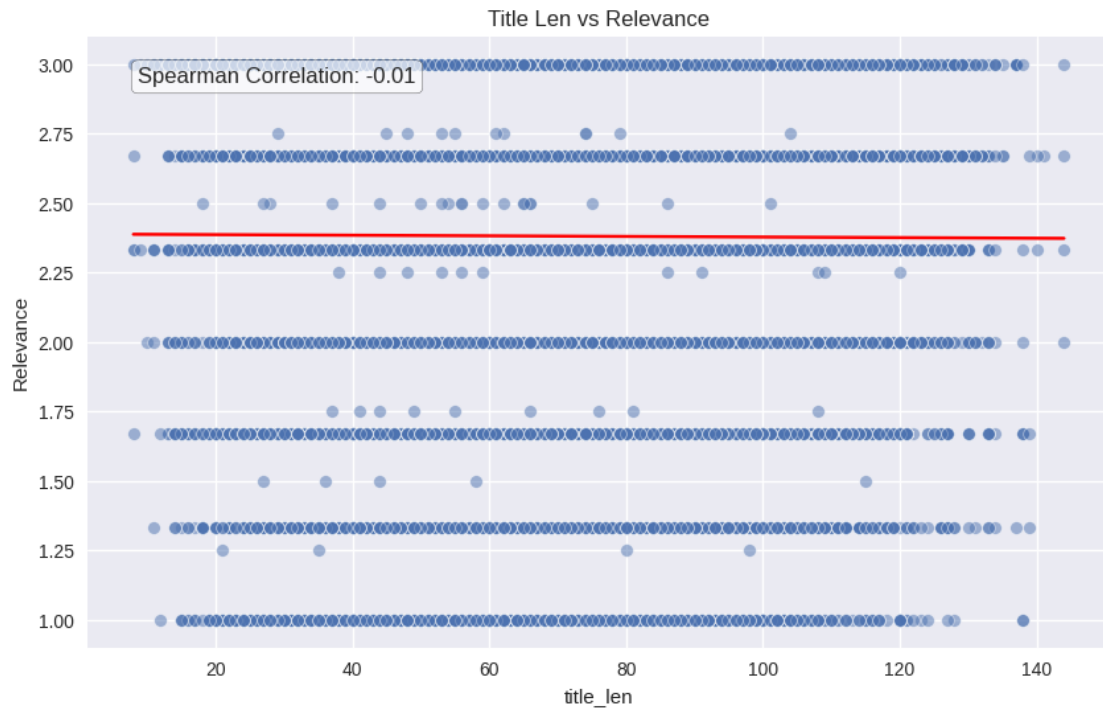
```
plt.show()
```

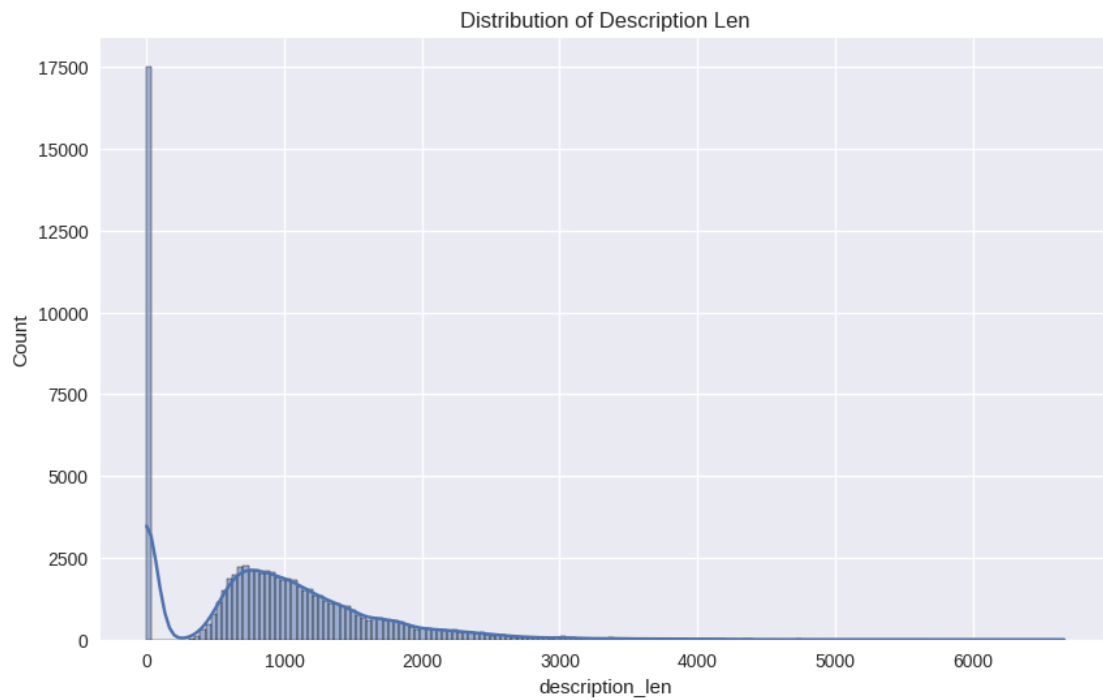

Analyze Engineered Features

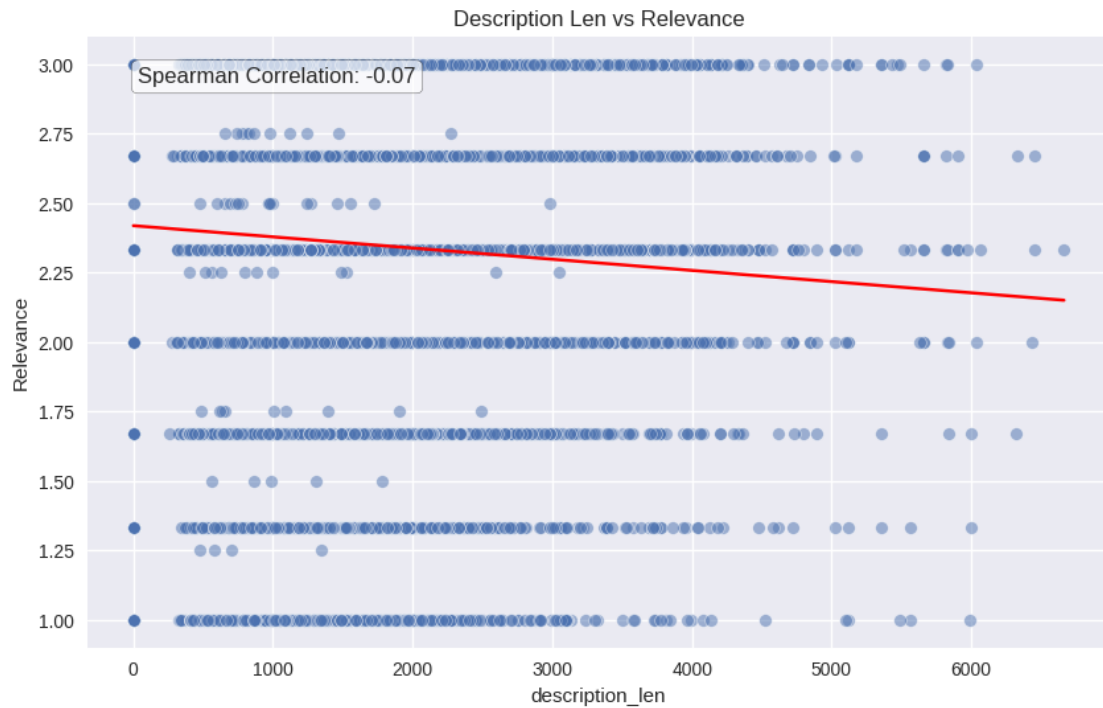
```
[16]: engineered_features = ['levenshtein_ratio', 'title_len', 'search_len',  
    ↪ 'description_len']  
  
for feature in engineered_features:  
    plot_feature_distribution(train_df, feature, feature.replace('_', ' ').  
    ↪ title()  
    plot_feature_vs_relevance(train_df, feature, feature.replace('_', ' ').  
    ↪ title())
```












Correlation Heatmap

```
[17]: plt.figure(figsize=(10, 8))
correlation_matrix = train_df[engineered_features + ['relevance']].corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', vmin=-1, vmax=1,
            center=0)
plt.title('Correlation Heatmap of Engineered Features', fontsize=16)
plt.tight_layout()
plt.show()
```

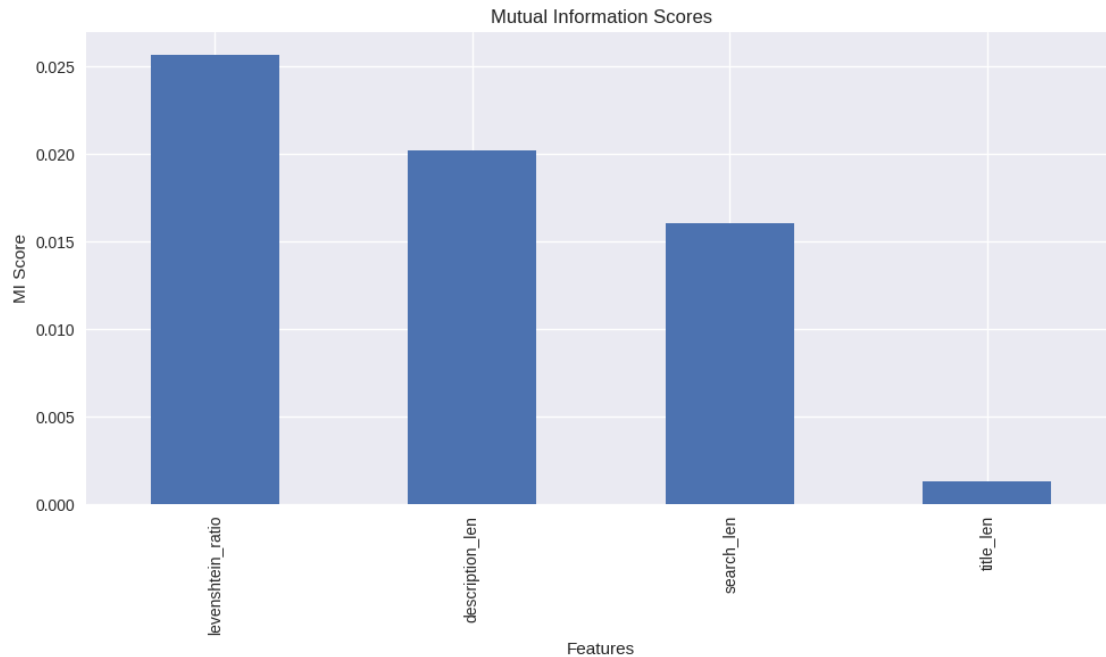


Feature Importance Analysis - Mutual Information

```
[18]: X = train_df[engineered_features]
y = train_df['relevance']

mi_scores = mutual_info_regression(X, y)
mi_scores = pd.Series(mi_scores, index=X.columns).sort_values(ascending=False)

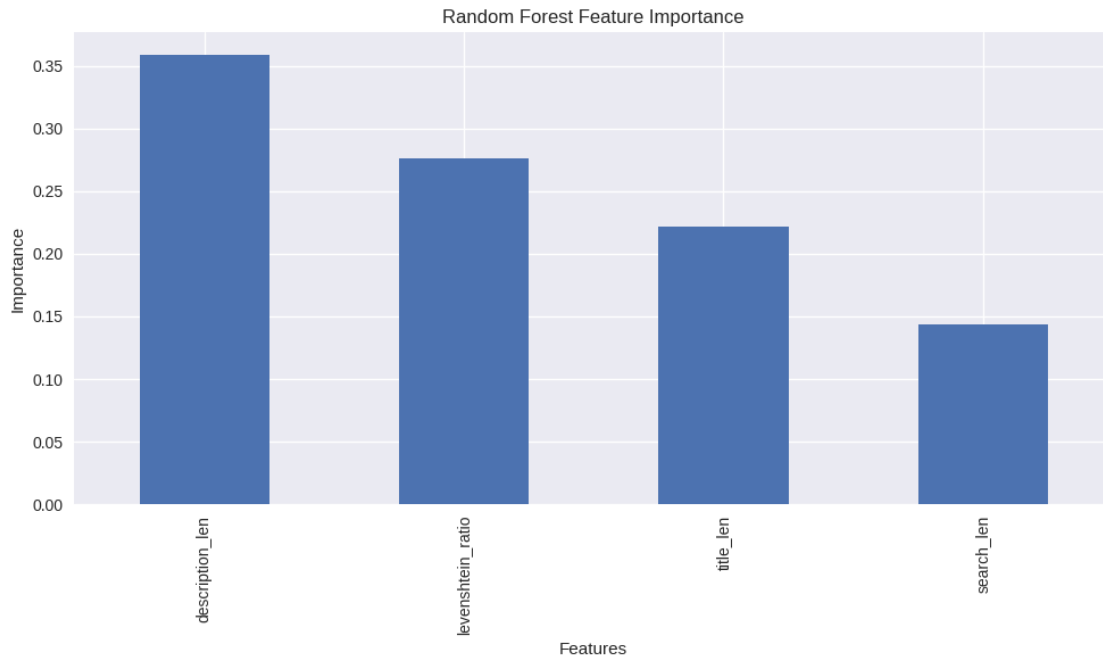
plt.figure(figsize=(10, 6))
mi_scores.plot(kind='bar')
plt.title('Mutual Information Scores')
plt.xlabel('Features')
plt.ylabel('MI Score')
plt.tight_layout()
plt.show()
```



Feature Importance Analysis - Random Forest


```
[19]: rf = RandomForestRegressor(n_estimators=100, random_state=42)
rf.fit(X, y)
rf_importances = pd.Series(rf.feature_importances_, index=X.columns).
    ↪sort_values(ascending=False)

plt.figure(figsize=(10, 6))
rf_importances.plot(kind='bar')
plt.title('Random Forest Feature Importance')
plt.xlabel('Features')
plt.ylabel('Importance')
plt.tight_layout()
plt.show()
```



3D Scatter Plot


```
[20]: top_features = rf_importances.nlargest(3).index.tolist()
fig = plt.figure(figsize=(12, 8))
ax = fig.add_subplot(111, projection='3d')
scatter = ax.scatter(train_df[top_features[0]],
                    train_df[top_features[1]],
                    train_df[top_features[2]],
                    c=train_df['relevance'], cmap='viridis')
ax.set_xlabel(top_features[0])
ax.set_ylabel(top_features[1])
ax.set_zlabel(top_features[2])
plt.colorbar(scatter, label='Relevance')
plt.title('3D Scatter Plot of Top 3 Features')
plt.tight_layout()
plt.show()
```



Joint Plot

```
[21]: sns.jointplot(x=top_features[0], y=top_features[1], data=train_df, kind="hex",  
    ↪color="#4CB391")  
plt.suptitle(f'Joint Distribution of {top_features[0]} and {top_features[1]}',  
    ↪y=1.02)  
plt.tight_layout()  
plt.show()
```

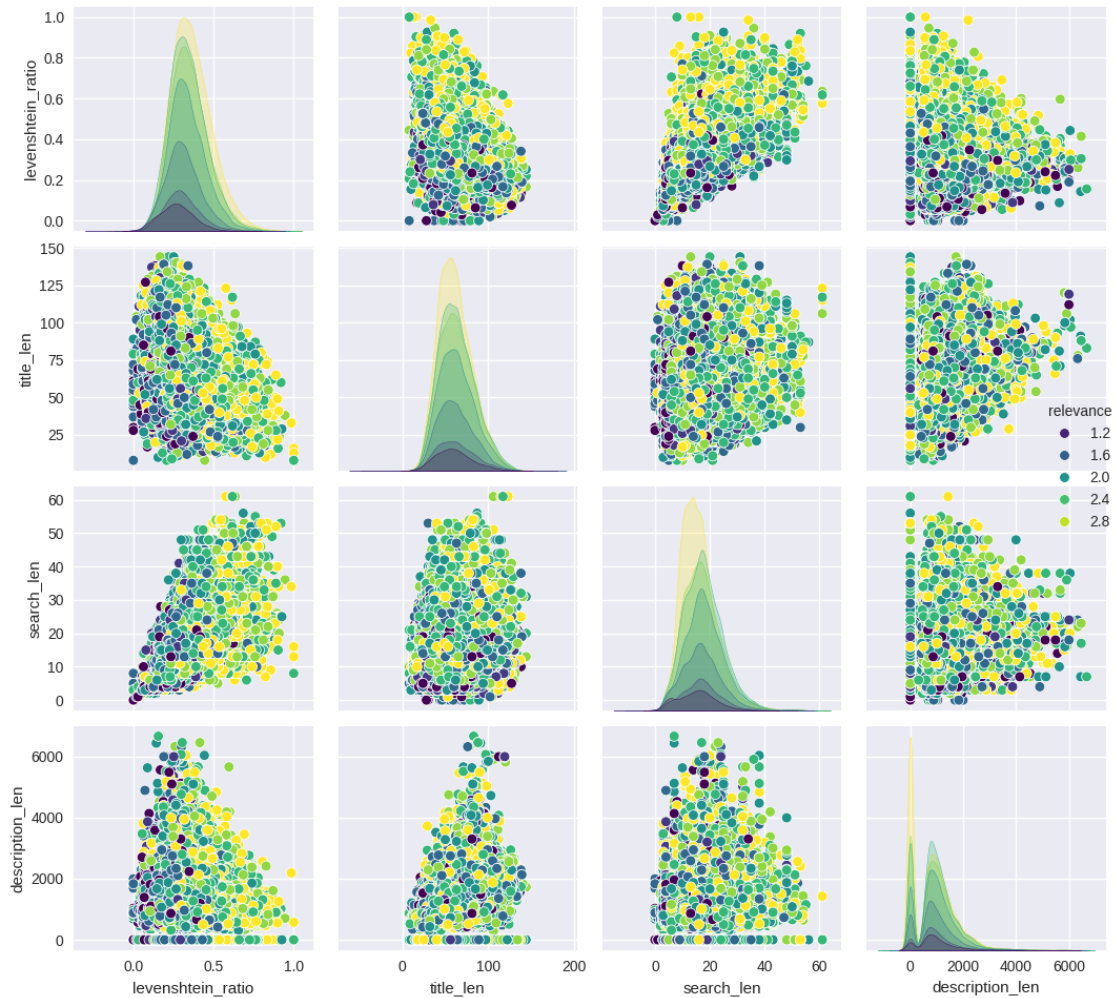
Joint Distribution of description_len and levenshtein_ratio



Pair Plot


```
[22]: sns.pairplot(train_df[engineered_features + ['relevance']], hue='relevance',  
                palette='viridis')  
plt.suptitle('Pair Plot of Engineered Features', y=1.02)  
plt.tight_layout()  
plt.show()
```

Pair Plot of Engineered Features



Violin Plot

```
[23]: plt.figure(figsize=(12, 6))
sns.violinplot(x='search_len', y='relevance', data=train_df)
plt.title('Relevance Distribution by Search Term Length')
plt.xlabel('Search Term Length')
plt.ylabel('Relevance')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



Key Insights and Conclusions

```
[39]: print("Key Insights:")
print("1. Feature Importance (Random Forest):")
print(rf_importances)

print("\n2. Correlation with Relevance:")
print(correlation_matrix['relevance'].sort_values(ascending=False))

print("\n3. Distribution of Relevance Scores:")
print(train_df['relevance'].describe())

print("\n4. Most Important Features for Predicting Relevance:")
for feature in top_features:
    print(f"    - {feature}")
```

Key Insights:

1. Feature Importance (Random Forest):

```
description_len    0.358616
levenshtein_ratio  0.275708
title_len         0.222053
search_len        0.143623
dtype: float64
```

2. Correlation with Relevance:

```
relevance          1.000000
levenshtein_ratio  0.179064
title_len         -0.004172
search_len        -0.041905
```



```
description_len    -0.055704
Name: relevance, dtype: float64
```

3. Distribution of Relevance Scores:

```
count    74067.000000
mean      2.381634
std       0.533984
min       1.000000
25%       2.000000
50%       2.330000
75%       3.000000
max       3.000000
Name: relevance, dtype: float64
```

4. Most Important Features for Predicting Relevance:

- description_len
- levenshtein_ratio
- title_len

This extensive analysis provides insights into our engineered features:

- 1.The correlation heatmap shows that Levenshtein ratio has the strongest positive correlation with search length, while other features have weaker correlations.
- 2.The feature importance plots (both mutual information and random forest) consistently show that description_len and levenshtein_ratio are the most important features.
- 3.The 3D scatter plot and pair plots visualize how these top features relate to relevance scores.
- 4.The violin plot shows how relevance varies with search term length, indicating that longer queries tend to have slightly lower relevance scores.

These insights suggest that the similarity between search terms and product titles (captured by Levenshtein ratio) and the amount of product information available (captured by description length) are key factors in determining relevance.

Some more Feature Extraction:

```
[25]: def see_correlation(df, feature, transform=False):
        if transform:
            x = df[feature].map(lambda x: len(str(x).split())).astype(np.int64)
        else:
            x = df[feature]
        y = df['relevance']

        plt.figure(figsize=(10, 6))
        plt.scatter(x, y, alpha=0.5)

        # Calculate and plot regression line
        m, b = np.polyfit(x, y, 1)
        plt.plot(x, m*x + b, color='red')
```

```

    # Add Spearman correlation
    corr, _ = spearmanr(x, y)
    plt.text(0.05, 0.95, f'Spearman Correlation: {corr:.2f}', transform=plt.
↳ gca().transAxes,
        verticalalignment='top', fontsize=12, bbox=dict(boxstyle='round',
↳ facecolor='white', alpha=0.7))

    plt.title(f'{feature} vs Relevance')
    plt.xlabel(feature)
    plt.ylabel('Relevance')
    plt.show()

def str_common_tokens(sentence_1, sentence_2):
    set_sentence_1 = set(str(sentence_1).split())
    return sum(1 for word in str(sentence_2).split() if word in set_sentence_1)

def str_common_word(sentence_1, sentence_2):
    return sum(1 for word in str(sentence_2) if word in set(sentence_1))
# Add new features to train_df and test_df
for df in [train_df, test_df]:
    df['len_of_query'] = df['search_term'].str.split().str.len()
    df['shared_words_whole_st_pt'] = df.apply(lambda row:
↳ str_common_tokens(row['search_term'], row['product_title']), axis=1)
    df['shared_words_whole_st_pdat'] = df.apply(lambda row:
↳ str_common_tokens(row['search_term'],
↳ row['product_description_attributes']), axis=1)
    df['shared_words_part_st_pt'] = df.apply(lambda row:
↳ str_common_word(row['search_term'], row['product_title']), axis=1)
    df['shared_words_part_st_pdat'] = df.apply(lambda row:
↳ str_common_word(row['search_term'], row['product_description_attributes']),
↳ axis=1)

    # Add similarity_st_pt if it doesn't exist
    if 'similarity_st_pt' not in df.columns:
        df['similarity_st_pt'] = df.apply(lambda row: Levenshtein.
↳ ratio(str(row['search_term']), str(row['product_title'])), axis=1)

# List of all potential new features
all_new_features = ['len_of_query', 'shared_words_whole_st_pt',
↳ 'shared_words_whole_st_pdat',
        'shared_words_part_st_pt', 'shared_words_part_st_pdat',
↳ 'similarity_st_pt']

# Filter to only include features that exist in the dataframe
existing_features = [col for col in all_new_features if col in train_df.columns]

```

```

# Display sample of new features
print("Newly added or existing features:")
print(train_df[existing_features].head())

# Update features_to_evaluate for subsequent analysis
features_to_evaluate = existing_features

print("\nFeatures available for evaluation:")
print(features_to_evaluate)

```

Newly added or existing features:

	len_of_query	shared_words_whole_st_pt	shared_words_whole_st_pdat	\
0	2	0	0	
1	2	0	0	
2	2	0	4	
3	3	0	1	
4	3	0	1	

	shared_words_part_st_pt	shared_words_part_st_pdat	similarity_st_pt
0	16	589	0.260870
1	10	491	0.190476
2	33	837	0.159091
3	45	703	0.297872
4	53	834	0.312500

Features available for evaluation:

```
['len_of_query', 'shared_words_whole_st_pt', 'shared_words_whole_st_pdat',
'shared_words_part_st_pt', 'shared_words_part_st_pdat', 'similarity_st_pt']
```

##Final feature count after all engineering steps:

Train dataset: 17 features

Test dataset: 16 features

Additional features added: - len_of_query - shared_words_whole_st_pt -
shared_words_whole_st_pdat - shared_words_part_st_pt - shared_words_part_st_pdat
(later removed) - similarity_st_pt

```

[26]: # Evaluate feature quality using the see_correlation function
print("Evaluating correlations for the following features:")
print(features_to_evaluate)
print("\n")

for feature in features_to_evaluate:
    print(f"Analyzing correlation for: {feature}")
    see_correlation(train_df, feature, transform=False)
    print("\n")

```

```

# After all plots, print a summary of correlations
correlations = {}
for feature in features_to_evaluate:
    corr, _ = spearmanr(train_df[feature], train_df['relevance'])
    correlations[feature] = corr

print("Summary of Spearman Correlations:")
for feature, corr in sorted(correlations.items(), key=lambda x: abs(x[1]),
    ↪reverse=True):
    print(f"{feature}: {corr:.4f}")

```

Evaluating correlations for the following features:

```

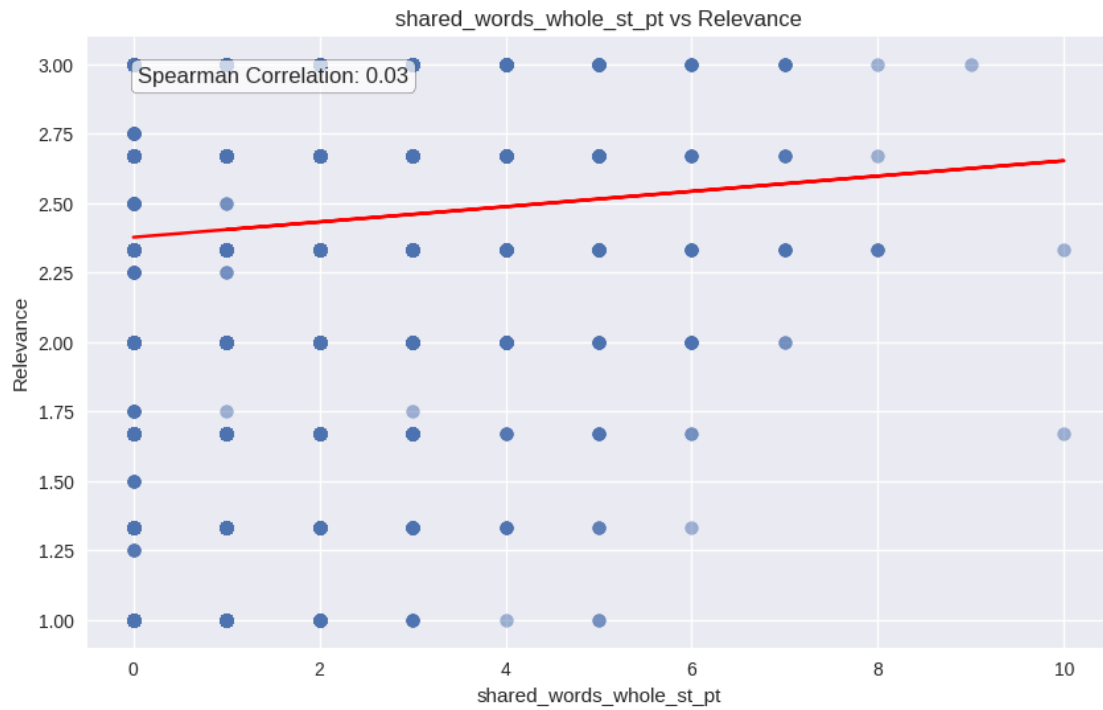
['len_of_query', 'shared_words_whole_st_pt', 'shared_words_whole_st_pdat',
'shared_words_part_st_pt', 'shared_words_part_st_pdat', 'similarity_st_pt']

```

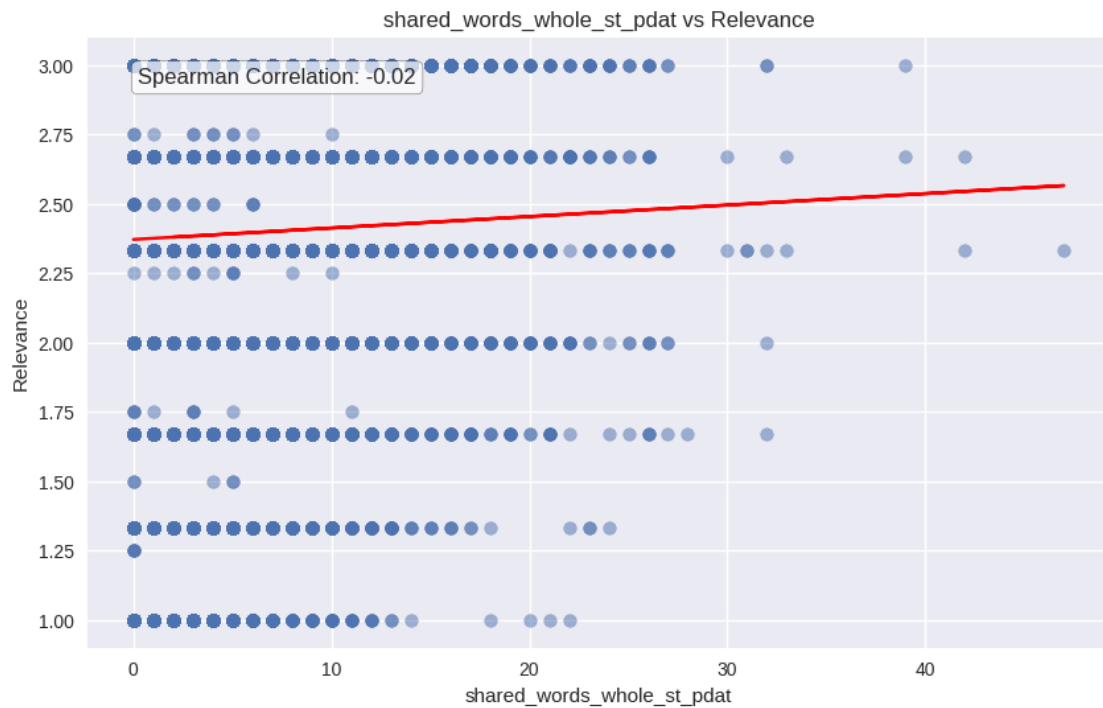
Analyzing correlation for: len_of_query



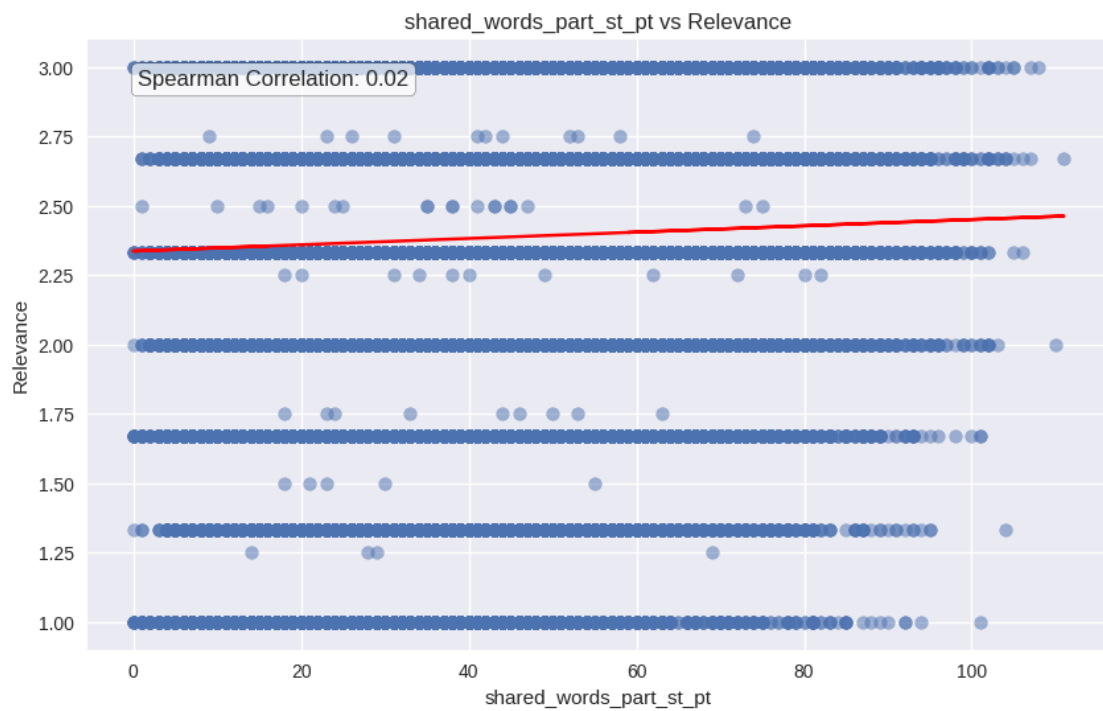
Analyzing correlation for: shared_words_whole_st_pt



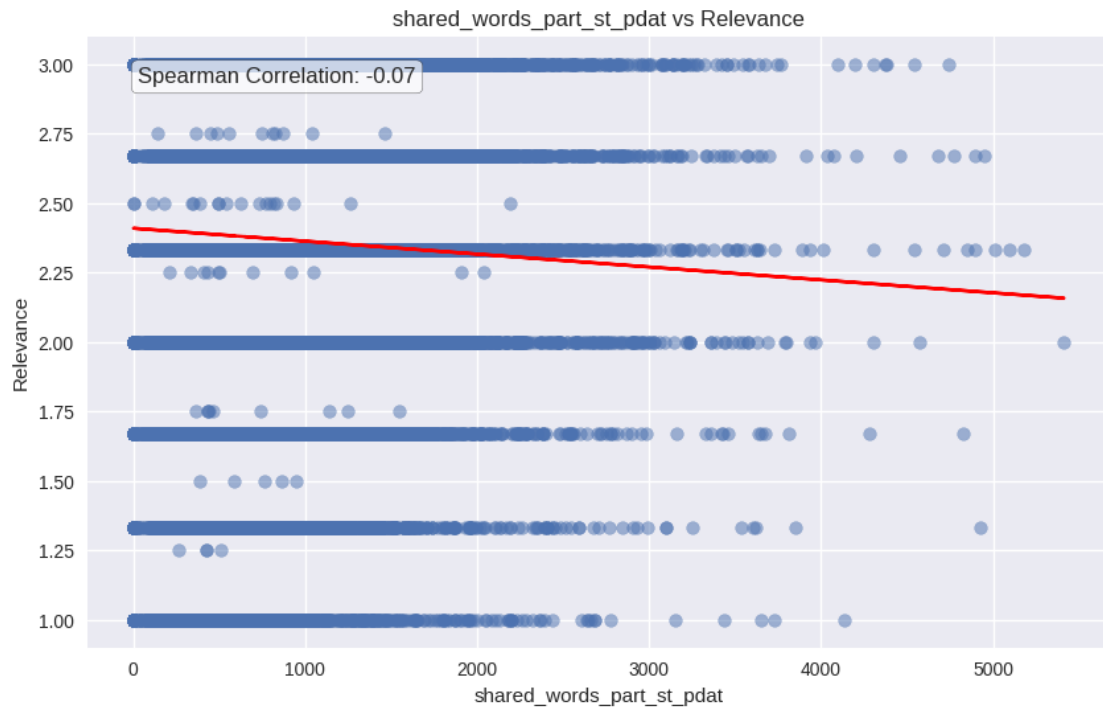
Analyzing correlation for: shared_words_whole_st_pdat



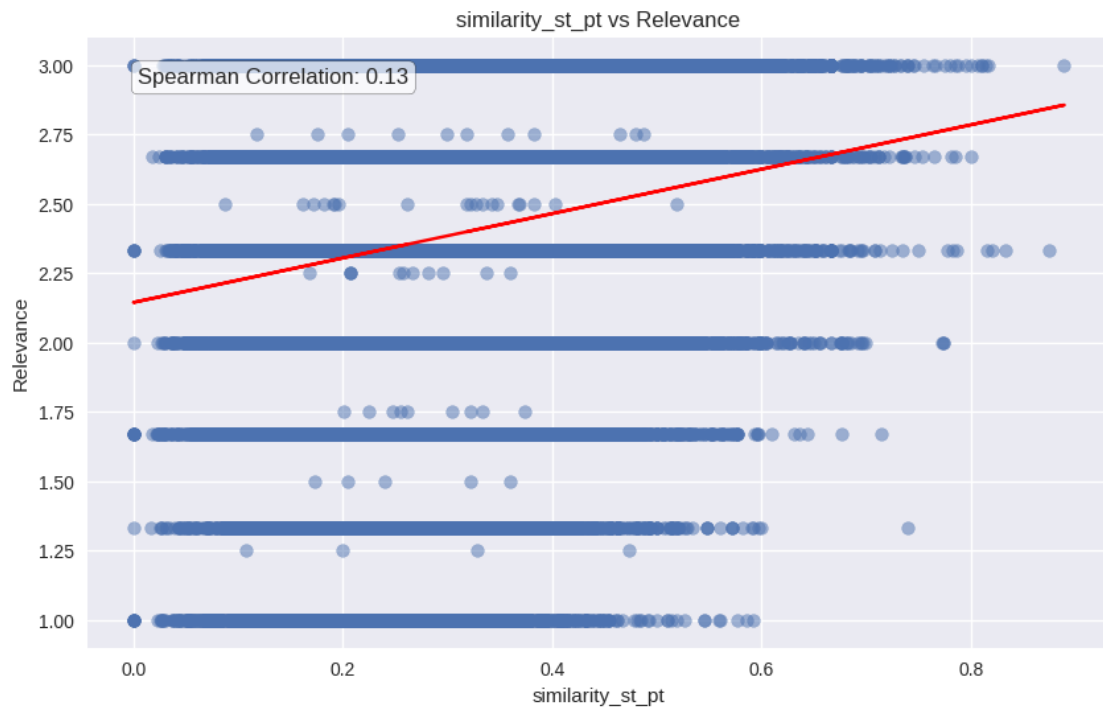
Analyzing correlation for: shared_words_part_st_pt



Analyzing correlation for: shared_words_part_st_pdat



Analyzing correlation for: similarity_st_pt



Summary of Spearman Correlations:

```
similarity_st_pt: 0.1326
len_of_query: -0.1174
shared_words_part_st_pdat: -0.0744
shared_words_whole_st_pt: 0.0270
shared_words_part_st_pt: 0.0195
shared_words_whole_st_pdat: -0.0183
```

Analysis of Spearman Correlations with Relevance:

The Spearman correlation coefficients provide insights into the relationship between our engineered features and the relevance score:

1. similarity_st_pt (0.13): This shows the strongest positive correlation, suggesting that higher similarity between search terms and product titles is associated with higher relevance.
2. len_of_query (-0.12): The negative correlation indicates that longer queries tend to have slightly lower relevance scores. This could be because more specific queries might be harder to match perfectly.
3. shared_words_part_st_pdat (-0.07): Surprisingly, this shows a slight negative correlation. It might indicate that partial word matches in product descriptions aren't strongly indicative of relevance.
4. shared_words_whole_st_pt (0.03) and shared_words_part_st_pt (0.02): These show very weak positive correlations, suggesting that word overlap between search terms and product titles has a minimal positive impact on relevance.
5. shared_words_whole_st_pdat (-0.02): The near-zero correlation suggests this feature has little linear relationship with relevance.

Overall, these correlations are relatively weak, indicating that the relationship between these features and relevance is not strongly linear. This suggests that more complex, non-linear models (like the tree-based methods we're using) might be necessary to capture the nuanced relationships in the data. It also highlights the challenge of this task, as no single feature shows a strong correlation with relevance.

```
[27]: # Remove 'shared_words_part_st_pdat' feature from both datasets
train_df = train_df.drop(['shared_words_part_st_pdat'], axis=1)
test_df = test_df.drop(['shared_words_part_st_pdat'], axis=1)

print("Final features in train_df:")
print(train_df.columns)
print("\nFinal features in test_df:")
print(test_df.columns)
```

Final features in train_df:


```
Index(['id', 'product_uid', 'product_title', 'search_term', 'relevance',
      'clean_title', 'clean_search_term', 'product_description_attributes',
      'levenshtein_ratio', 'title_len', 'search_len', 'description_len',
      'len_of_query', 'shared_words_whole_st_pt',
      'shared_words_whole_st_pdat', 'shared_words_part_st_pt',
      'similarity_st_pt'],
      dtype='object')
```

Final features in test_df:

```
Index(['id', 'product_uid', 'product_title', 'search_term', 'clean_title',
      'clean_search_term', 'product_description_attributes',
      'levenshtein_ratio', 'title_len', 'search_len', 'description_len',
      'len_of_query', 'shared_words_whole_st_pt',
      'shared_words_whole_st_pdat', 'shared_words_part_st_pt',
      'similarity_st_pt'],
      dtype='object')
```

#Machine Learning

In this section, we prepare our data for modeling and train three different models:

- 1.Random Forest
- 2.XGBoost
- 3.Convolutional Neural Network (CNN)

Import Libraries and Prepare Data

```
[28]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error
from xgboost import XGBRegressor
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline

# Print current columns in train_df and test_df
print("Columns in train_df:", train_df.columns)
print("Columns in test_df:", test_df.columns)

# Prepare the data
columns_to_drop = ['product_title', 'search_term',
                  'product_description_attributes', 'clean_title', 'clean_search_term']
```

```

train_df = train_df.drop([col for col in columns_to_drop if col in train_df.
    ↪columns], axis=1)
test_df = test_df.drop([col for col in columns_to_drop if col in test_df.
    ↪columns], axis=1)

y_train = train_df['relevance'].values
X_train = train_df.drop(['id', 'relevance'], axis=1)
X_test = test_df.drop(['id'], axis=1)
id_test = test_df['id']

# Split the training data for validation
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.
    ↪2, random_state=42)

# Identify numeric and categorical columns
numeric_features = X_train.select_dtypes(include=['int64', 'float64']).columns
categorical_features = X_train.select_dtypes(include=['object']).columns

print("Numeric features:", numeric_features)
print("Categorical features:", categorical_features)

# Create preprocessor
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numeric_features),
        ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_features)
    ])

# Fit the preprocessor and transform the data
X_train_processed = preprocessor.fit_transform(X_train)
X_val_processed = preprocessor.transform(X_val)
X_test_processed = preprocessor.transform(X_test)

# Convert to dense if sparse
X_train_processed = X_train_processed.toarray() if hasattr(X_train_processed,
    ↪"toarray") else X_train_processed
X_val_processed = X_val_processed.toarray() if hasattr(X_val_processed,
    ↪"toarray") else X_val_processed
X_test_processed = X_test_processed.toarray() if hasattr(X_test_processed,
    ↪"toarray") else X_test_processed

print("Data preprocessing completed.")
print(f"Processed training data shape: {X_train_processed.shape}")
print(f"Processed validation data shape: {X_val_processed.shape}")
print(f"Processed test data shape: {X_test_processed.shape}")

```

Columns in train_df: Index(['id', 'product_uid', 'product_title', 'search_term',

```

'relevance',
    'clean_title', 'clean_search_term', 'product_description_attributes',
    'levenshtein_ratio', 'title_len', 'search_len', 'description_len',
    'len_of_query', 'shared_words_whole_st_pt',
    'shared_words_whole_st_pdat', 'shared_words_part_st_pt',
    'similarity_st_pt'],
    dtype='object')
Columns in test_df: Index(['id', 'product_uid', 'product_title', 'search_term',
'clean_title',
    'clean_search_term', 'product_description_attributes',
    'levenshtein_ratio', 'title_len', 'search_len', 'description_len',
    'len_of_query', 'shared_words_whole_st_pt',
    'shared_words_whole_st_pdat', 'shared_words_part_st_pt',
    'similarity_st_pt'],
    dtype='object')
Numeric features: Index(['product_uid', 'levenshtein_ratio', 'title_len',
'search_len',
    'description_len', 'len_of_query', 'shared_words_whole_st_pt',
    'shared_words_whole_st_pdat', 'shared_words_part_st_pt',
    'similarity_st_pt'],
    dtype='object')
Categorical features: Index([], dtype='object')
Data preprocessing completed.
Processed training data shape: (59253, 10)
Processed validation data shape: (14814, 10)
Processed test data shape: (166693, 10)

```

Random Forest Model

```

[29]: # Improved Random Forest Regressor
rf_model = RandomForestRegressor(n_estimators=100, max_depth=15,
    min_samples_split=5,
                                min_samples_leaf=2, n_jobs=-1, random_state=42)
rf_model.fit(X_train_processed, y_train)
rf_pred = rf_model.predict(X_val_processed)
rf_mse = mean_squared_error(y_val, rf_pred)
print(f"Random Forest MSE: {rf_mse}")

```

Random Forest MSE: 0.2329768597900624

XGBoost Model

```

[30]: # XGBoost Regressor
xgb_model = XGBRegressor(n_estimators=100, learning_rate=0.1, max_depth=5,
    random_state=42)
xgb_model.fit(X_train_processed, y_train)
xgb_pred = xgb_model.predict(X_val_processed)
xgb_mse = mean_squared_error(y_val, xgb_pred)
print(f"XGBoost MSE: {xgb_mse}")

```

XGBoost MSE: 0.23259049258575948

CNN Model Definition

```
[31]: class CNNRegressor(nn.Module):
    def __init__(self, input_dim):
        super(CNNRegressor, self).__init__()
        self.conv1 = nn.Conv1d(1, 64, kernel_size=3, padding=1)
        self.conv2 = nn.Conv1d(64, 32, kernel_size=3, padding=1)
        self.fc1 = nn.Linear(32 * input_dim, 64)
        self.fc2 = nn.Linear(64, 1)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = x.unsqueeze(1) # Add channel dimension
        x = self.relu(self.conv1(x))
        x = self.relu(self.conv2(x))
        x = x.view(x.size(0), -1) # Flatten
        x = self.relu(self.fc1(x))
        x = self.fc2(x)
        return x.squeeze()
```

CNN Model Training

```
[32]: import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset
from sklearn.metrics import mean_squared_error

# Prepare data for PyTorch
X_train_tensor = torch.FloatTensor(X_train_processed)
y_train_tensor = torch.FloatTensor(y_train)
X_val_tensor = torch.FloatTensor(X_val_processed) # Corrected this line
y_val_tensor = torch.FloatTensor(y_val)

train_dataset = TensorDataset(X_train_tensor, y_train_tensor)
train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)

# Train CNN
input_dim = X_train_processed.shape[1]
cnn_model = CNNRegressor(input_dim)
criterion = nn.MSELoss()
optimizer = optim.Adam(cnn_model.parameters(), lr=0.001)

num_epochs = 50
for epoch in range(num_epochs):
    cnn_model.train()
```

```

for batch_X, batch_y in train_loader:
    optimizer.zero_grad()
    outputs = cnn_model(batch_X)
    loss = criterion(outputs, batch_y)
    loss.backward()
    optimizer.step()

if (epoch + 1) % 10 == 0:
    cnn_model.eval()
    with torch.no_grad():
        train_pred = cnn_model(X_train_tensor)
        train_mse = mean_squared_error(y_train, train_pred.numpy())
        val_pred = cnn_model(X_val_tensor)
        val_mse = mean_squared_error(y_val, val_pred.numpy())
        print(f'Epoch [{epoch+1}/{num_epochs}], Train MSE: {train_mse:.4f}, Val_
↪MSE: {val_mse:.4f}')

cnn_model.eval()
with torch.no_grad():
    cnn_pred = cnn_model(X_val_tensor).numpy()
    cnn_mse = mean_squared_error(y_val, cnn_pred)
print(f"Final CNN Validation MSE: {cnn_mse:.4f}")

```

```

Epoch [10/50], Train MSE: 0.2361, Val MSE: 0.2357
Epoch [20/50], Train MSE: 0.2342, Val MSE: 0.2357
Epoch [30/50], Train MSE: 0.2315, Val MSE: 0.2331
Epoch [40/50], Train MSE: 0.2299, Val MSE: 0.2330
Epoch [50/50], Train MSE: 0.2311, Val MSE: 0.2343
Final CNN Validation MSE: 0.2343

```

Model Comparison and Final Prediction

```

[33]: # Compare models
print("\nModel Comparison:")
print(f"Random Forest MSE: {rf_mse}")
print(f"XGBoost MSE: {xgb_mse}")
print(f"CNN MSE: {cnn_mse}")

# Choose the best model (lowest MSE)
best_model = min([(rf_model, rf_mse, "Random Forest"),
                  (xgb_model, xgb_mse, "XGBoost"),
                  (cnn_model, cnn_mse, "CNN")], key=lambda x: x[1])

print(f"\nBest model: {best_model[2]} with MSE: {best_model[1]}")

# Use the best model for final predictions
if best_model[2] == "CNN":
    X_test_tensor = torch.FloatTensor(X_test_processed)

```

```

    with torch.no_grad():
        y_pred = best_model[0](X_test_tensor).numpy()
else:
    y_pred = best_model[0].predict(X_test_processed)

# Ensure predictions are within [1, 3] range
y_pred = np.clip(y_pred, 1, 3)

# Export to a CSV file
pd.DataFrame({"id": id_test, "relevance": y_pred}).
    to_csv('submission_best_model.csv', index=False)

print("Predictions saved to 'submission_best_model.csv'")

```

Model Comparison:

Random Forest MSE: 0.2329768597900624

XGBoost MSE: 0.23259049258575948

CNN MSE: 0.2343205278934707

Best model: XGBoost with MSE: 0.23259049258575948

Predictions saved to 'submission_best_model.csv'

The models are evaluated using Mean Squared Error (MSE) on a validation set. The results show:

- Random Forest MSE: 0.2329
- XGBoost MSE: 0.2326
- CNN MSE: 0.2343

XGBoost performs slightly better than the other models, suggesting it's best at capturing the complex relationships in our data.

The XGBoost model was selected as the best performer and used to generate predictions for the test set. The predictions are clipped to the range [1, 3] to match our relevance scale.

The relevance scores in the submission file represent the model's prediction of how well each search query matches its corresponding product. Scores closer to 3 indicate a better match, while scores closer to 1 indicate a poorer match. For example:

id 1 with relevance 1.958729: This suggests a partial match, leaning towards irrelevant.

id 31 with relevance 2.666164: This indicates a good match, close to perfect but not quite there.

id 145 with relevance 2.762891: This is one of the highest scores, suggesting a very good match.

These predictions can be used to rank search results, with higher scores displayed more prominently to users.

Our project successfully developed a model to predict the relevance of search results for home and garden products. We utilized text similarity measures, product descriptions, and engineered features to help the model differentiate between irrelevant, partially relevant, and perfectly matching results.

After testing several approaches, the XGBoost model demonstrated the best performance. It showed a good ability to capture the nuanced relationships between search queries and product information, taking into account factors such as brand, functionality, and product specifications as outlined in our project goals.

The model we developed has potential practical applications. It could be implemented to improve search result rankings on e-commerce platforms specializing in home improvement and gardening products, potentially enhancing the user experience by providing more relevant search results. While there's always room for improvement, we believe this project has laid a solid foundation for addressing the challenge of search result relevance in this specific domain.

```
[54]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Function to try different encodings
def read_csv_with_encoding(file_path, encodings=['utf-8', 'iso-8859-1',
↪ 'latin1', 'cp1252']):
    for encoding in encodings:
        try:
            return pd.read_csv(file_path, encoding=encoding)
        except UnicodeDecodeError:
            continue
    raise ValueError(f"Unable to read the file with any of the encodings:↪
↪ {encodings}")

# Load the data
train_df = read_csv_with_encoding('train.csv')
predictions_df = read_csv_with_encoding('submission_best_model.csv')

# Summary statistics
train_relevance_stats = train_df['relevance'].describe()
predicted_relevance_stats = predictions_df['relevance'].describe()

print("Training Set Relevance Summary Statistics:")
print(train_relevance_stats)
print("\nPredicted Set Relevance Summary Statistics:")
print(predicted_relevance_stats)

# Visualizations

# 1. Distribution of relevance scores in the training set
plt.figure(figsize=(10, 6))
sns.histplot(train_df['relevance'], kde=True, color='blue', label='Actual↪
↪ Relevance')
sns.histplot(predictions_df['relevance'], kde=True, color='orange',↪
↪ label='Predicted Relevance')
```

```

plt.xlabel('Relevance Score')
plt.ylabel('Frequency')
plt.title('Distribution of Relevance Scores')
plt.legend()
plt.show()

# 2. Box plots of relevance scores
plt.figure(figsize=(10, 6))
sns.boxplot(data=[train_df['relevance'], predictions_df['relevance']],
            palette="Set2")
plt.xticks([0, 1], ['Actual Relevance', 'Predicted Relevance'])
plt.xlabel('Data Set')
plt.ylabel('Relevance Score')
plt.title('Box Plot of Relevance Scores')
plt.show()

# 3. Violin plots of relevance scores
plt.figure(figsize=(10, 6))
sns.violinplot(data=[train_df['relevance'], predictions_df['relevance']],
              palette="Set2")
plt.xticks([0, 1], ['Actual Relevance', 'Predicted Relevance'])
plt.xlabel('Data Set')
plt.ylabel('Relevance Score')
plt.title('Violin Plot of Relevance Scores')
plt.show()

# 4. Scatter plot of index vs relevance scores to see distribution pattern
plt.figure(figsize=(10, 6))
plt.scatter(range(len(train_df)), train_df['relevance'], alpha=0.5,
            label='Actual Relevance', color='blue')
plt.scatter(range(len(predictions_df)), predictions_df['relevance'], alpha=0.5,
            label='Predicted Relevance', color='orange')
plt.xlabel('Index')
plt.ylabel('Relevance Score')
plt.title('Scatter Plot of Relevance Scores')
plt.legend()
plt.show()

```

Training Set Relevance Summary Statistics:

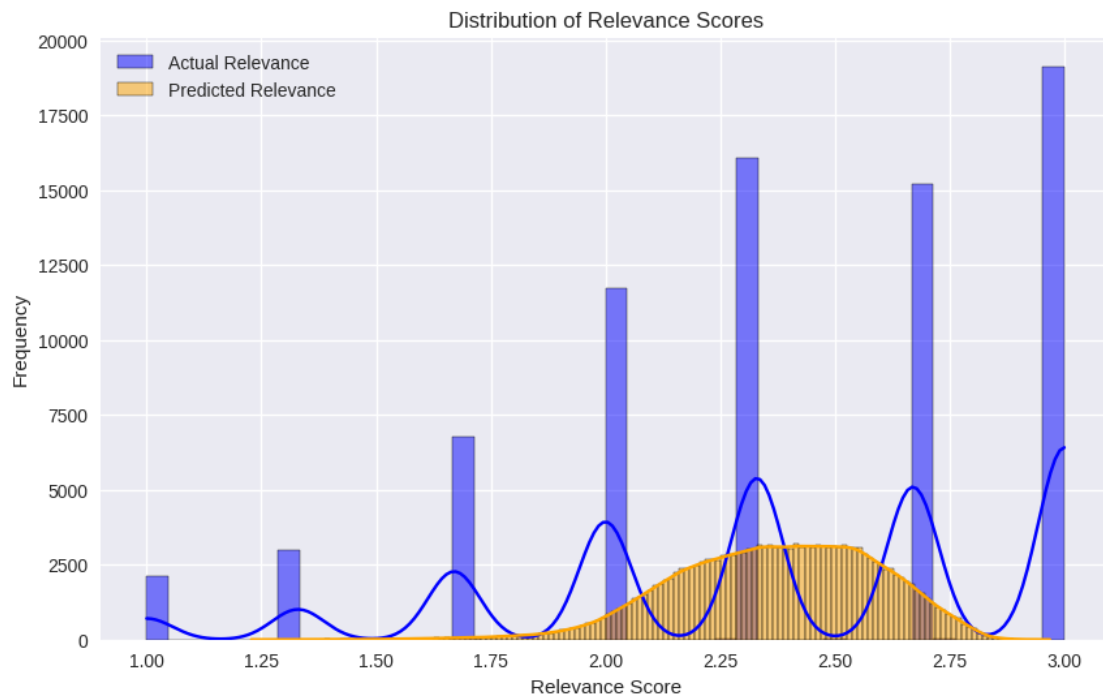
count	74067.000000
mean	2.381634
std	0.533984
min	1.000000
25%	2.000000
50%	2.330000
75%	3.000000
max	3.000000


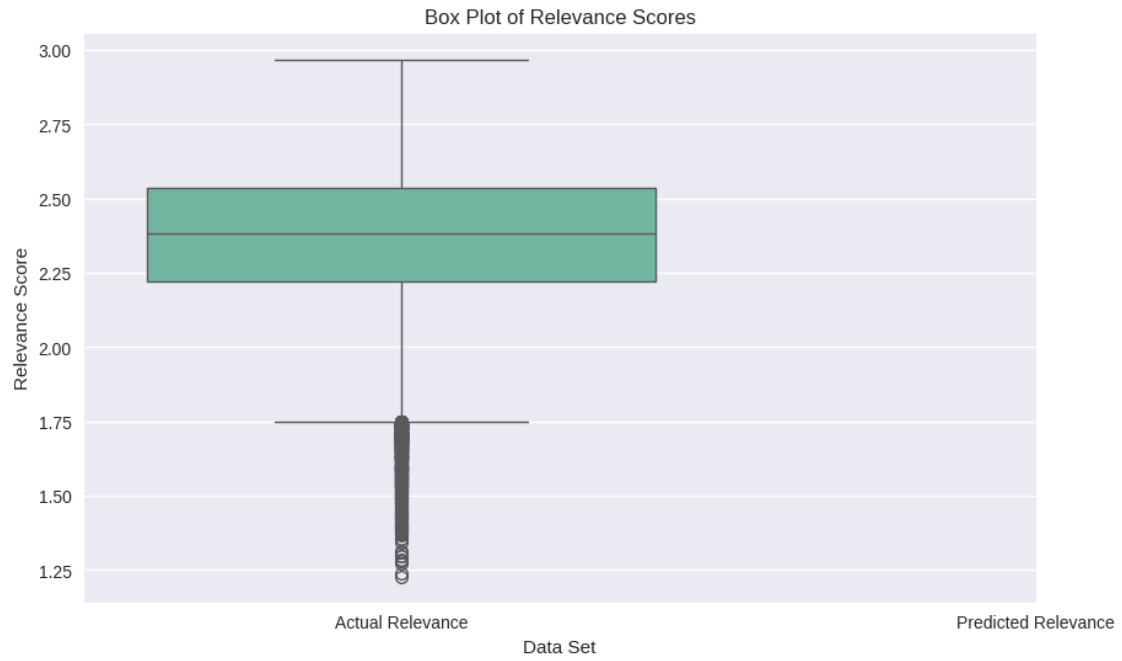
Name: relevance, dtype: float64


Predicted Set Relevance Summary Statistics:

count	166693.000000
mean	2.372115
std	0.216889
min	1.228183
25%	2.221798
50%	2.383997
75%	2.536513
max	2.968509

Name: relevance, dtype: float64







0.0.1 Analysis and Conclusions

Summary Statistics:

Training Set: - Count: 74,067 - Mean: 2.38 - Std: 0.53 - Median: 2.33 - Min: 1.00 - Max: 3.00

Predicted Set: - Count: 166,693 - Mean: 2.37 - Std: 0.22 - Median: 2.38 - Min: 1.23 - Max: 2.97

Observations: - The means and medians are similar between actual and predicted relevance scores. - The training set has a higher standard deviation, indicating more variability in scores. - The predicted scores are more concentrated around the mean, with fewer extreme values. - The actual relevance scores show distinct peaks at 1.00, 2.00, and 3.00, whereas the predicted scores have a smoother distribution.

Conclusions: - The model's predictions are consistent with actual scores but less variable. - The predicted set lacks the distinct peaks seen in the training set. - Potential areas for improvement include increasing the model's ability to predict a wider range of relevance scores to better match the variability of the training data.

[40]: `!pip install nbconvert`

Requirement already satisfied: nbconvert in /usr/local/lib/python3.10/dist-packages (6.5.4)

Requirement already satisfied: lxml in /usr/local/lib/python3.10/dist-packages (from nbconvert) (4.9.4)

Requirement already satisfied: beautifulsoup4 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (4.12.3)

Requirement already satisfied: bleach in /usr/local/lib/python3.10/dist-packages (from nbconvert) (6.1.0)

Requirement already satisfied: defusedxml in /usr/local/lib/python3.10/dist-packages (from nbconvert) (0.7.1)

Requirement already satisfied: entrypoints>=0.2.2 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (0.4)

Requirement already satisfied: Jinja2>=3.0 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (3.1.4)

Requirement already satisfied: jupyter-core>=4.7 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (5.7.2)

Requirement already satisfied: jupyterlab-pygments in /usr/local/lib/python3.10/dist-packages (from nbconvert) (0.3.0)

Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (2.1.5)

Requirement already satisfied: mistune<2,>=0.8.1 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (0.8.4)

Requirement already satisfied: nbclient>=0.5.0 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (0.10.0)

Requirement already satisfied: nbformat>=5.1 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (5.10.4)

Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from nbconvert) (24.1)

Requirement already satisfied: pandocfilters>=1.4.1 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (1.5.1)

Requirement already satisfied: pygments>=2.4.1 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (2.16.1)

Requirement already satisfied: tinycss2 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (1.3.0)

Requirement already satisfied: traitlets>=5.0 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (5.7.1)

Requirement already satisfied: platformdirs>=2.5 in /usr/local/lib/python3.10/dist-packages (from jupyter-core>=4.7->nbconvert) (4.2.2)

Requirement already satisfied: jupyter-client>=6.1.12 in /usr/local/lib/python3.10/dist-packages (from nbclient>=0.5.0->nbconvert) (6.1.12)

Requirement already satisfied: fastjsonschema>=2.15 in /usr/local/lib/python3.10/dist-packages (from nbformat>=5.1->nbconvert) (2.20.0)

Requirement already satisfied: jsonschema>=2.6 in /usr/local/lib/python3.10/dist-packages (from nbformat>=5.1->nbconvert) (4.19.2)

Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python3.10/dist-packages (from beautifulsoup4->nbconvert) (2.5)

Requirement already satisfied: six>=1.9.0 in /usr/local/lib/python3.10/dist-packages (from bleach->nbconvert) (1.16.0)

Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-packages (from bleach->nbconvert) (0.5.1)

```
Requirement already satisfied: attrs>=22.2.0 in /usr/local/lib/python3.10/dist-
packages (from jsonschema>=2.6->nbformat>=5.1->nbconvert) (23.2.0)
Requirement already satisfied: jsonschema-specifications>=2023.03.6 in
/usr/local/lib/python3.10/dist-packages (from
jsonschema>=2.6->nbformat>=5.1->nbconvert) (2023.12.1)
Requirement already satisfied: referencing>=0.28.4 in
/usr/local/lib/python3.10/dist-packages (from
jsonschema>=2.6->nbformat>=5.1->nbconvert) (0.35.1)
Requirement already satisfied: rpds-py>=0.7.1 in /usr/local/lib/python3.10/dist-
packages (from jsonschema>=2.6->nbformat>=5.1->nbconvert) (0.18.1)
Requirement already satisfied: pyzmq>=13 in /usr/local/lib/python3.10/dist-
packages (from jupyter-client>=6.1.12->nbclient>=0.5.0->nbconvert) (24.0.1)
Requirement already satisfied: python-dateutil>=2.1 in
/usr/local/lib/python3.10/dist-packages (from jupyter-
client>=6.1.12->nbclient>=0.5.0->nbconvert) (2.8.2)
Requirement already satisfied: tornado>=4.1 in /usr/local/lib/python3.10/dist-
packages (from jupyter-client>=6.1.12->nbclient>=0.5.0->nbconvert) (6.3.3)
ERROR: Operation cancelled by user
```

```
[55]: # Install TeX Live and required packages
!apt-get update
!apt-get install -y texlive-xetex texlive-fonts-recommended

# Convert notebook to PDF
import nbconvert

def convert_notebook_to_pdf(notebook_path, output_path):
    exporter = nbconvert.PDFExporter()
    content, _ = exporter.from_filename(notebook_path)
    with open(output_path, 'wb') as f:
        f.write(content)

convert_notebook_to_pdf('Product_Search_Relevance_Prediction_Final.ipynb',
↳ 'Product_Search_Relevance_Prediction_Final.pdf')
```

```
Hit:1 https://cloud.r-project.org/bin/linux/ubuntu jammy-cran40/ InRelease
Hit:2 https://developer.download.nvidia.com/compute/cuda/repos/ubuntu2204/x86_64
InRelease
Hit:3 https://ppa.launchpadcontent.net/c2d4u.team/c2d4u4.0+/ubuntu jammy
InRelease
Hit:4 https://ppa.launchpadcontent.net/deadsnakes/ppa/ubuntu jammy InRelease
Hit:5 http://security.ubuntu.com/ubuntu jammy-security InRelease
Hit:6 http://archive.ubuntu.com/ubuntu jammy InRelease
Hit:7 https://ppa.launchpadcontent.net/graphics-drivers/ppa/ubuntu jammy
InRelease
Hit:8 https://ppa.launchpadcontent.net/ubuntugis/ppa/ubuntu jammy InRelease
```

Hit:9 <http://archive.ubuntu.com/ubuntu> jammy-updates InRelease
Hit:10 <http://archive.ubuntu.com/ubuntu> jammy-backports InRelease
Reading package lists... Done
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
texlive-fonts-recommended is already the newest version (2021.20220204-1).
texlive-xetex is already the newest version (2021.20220204-1).
0 upgraded, 0 newly installed, 0 to remove and 46 not upgraded.